

# OpenMP Case Study: Bubble Sort

**Mike Bailey**

mjb@cs.oregonstate.edu

**Oregon State University**




	Step #					
original	0	1	2	3	4	5
6	5	5	3	3	1	1
5	6	3	5	1	3	2
4	3	6	1	5	2	3
3	4	1	6	2	5	4
2	1	4	2	6	4	5
1	2	2	4	4	6	6

## A Special Parallel Design Pattern

2

Implementing a Bubble Sort in parallel is an example of a special design pattern called ***Even-Odd***, or ***Red-Black***



# Non-threaded Bubble Sort

3

N = 6

```
#include <algorithm>
...
for( int i = 0; i < N; i++ )
{
    bool stop = true;

    for( int j = 0; j < N-1; j++ )
    {
        if( B[ j ] > B[ j+1 ] )
        {
            std::swap( B[ j ], B[ j+1 ] );
            stop = false;
        }
    }

    if( stop )
        break;
}
```

	Step #				
original	0	1	2	3	4
6	5	4	3	2	1
5	4	3	2	1	2
4	3	2	1	3	3
3	2	1	4	4	4
2	1	5	5	5	5
1	6	6	6	6	6

## Why Can't This Version of the Bubble Sort Be Run in Parallel?

4

```
#include <algorithm>
...
for( int i = 0; i < N; i++ )
{
    bool stop = true;

    if( B[ 0 ] > B[ 1 ] )
    {
        std::swap( B[ 0 ], B[ 1 ] );
        stop = false;
    }

    if( B[ 1 ] < B[ 2 ] )
    {
        std::swap( B[ 1 ], B[ 2 ] );
        stop = false;
    }

    if( B[ 2 ] > B[ 3 ] )
    {
        std::swap( B[ 2 ], B[ 3 ] );
        stop = false;
    }

    ...

    if( stop )
        break;
}
```

Let's unroll the inner (j) loop so we can see what the for-loop really looks like.

Suppose each of these if-blocks gets assigned to a different thread (remember that OpenMP tries to assign different for-loop passes to different threads).

Remembering that we have no explicit control over thread scheduling, notice that both the first and second if-blocks are both reading from and writing to **B[1]**. There is no synchronization to control in which order this is happening. We have a classic **Race Condition**.

The solution is, in one pass, allow a single thread access to B[0] and B[1], another thread access to B[2] and B[3], another thread access to B[4] and B[5], etc.

Then, in the next pass, allow a single thread access to B[1] and B[2], another thread access to B[3] and B[4], another thread access to B[5] and B[6], etc.

# Threaded Bubble Sort

5

```
#include <algorithm>
...
for( int i = 0; i < N; i++ )
{
    int first = i % 2;    // 0 if i is 0, 2, 4, ...
                        // 1 if i is 1, 3, 5, ...

    #pragma omp parallel for default(none),shared(A,first)

    for( int j = first; j < N-1; j += 2 )
    {
        if( A[ j ] > A[ j+1 ] )
        {
            std::swap( A[ j ], A[ j+1 ] );
        }
    }
}
```

N = 6

	Step #					
original	0	1	2	3	4	5
6	(5	5	(3	3	(1	1
5	(6	(3	(5	(1	(3	(2
4	(3	(6	(1	(5	(2	(3
3	(4	(1	(6	(2	(5	(4
2	(1	(4	(2	(6	(4	(5
1	(2	2	(4	4	(6	6

## A Comparison

6

$N = 6$

original	Step #				
	0	1	2	3	4
6	5	4	3	2	1
5	4	3	2	1	2
4	3	2	1	3	3
3	2	1	4	4	4
2	1	5	5	5	5
1	6	6	6	6	6

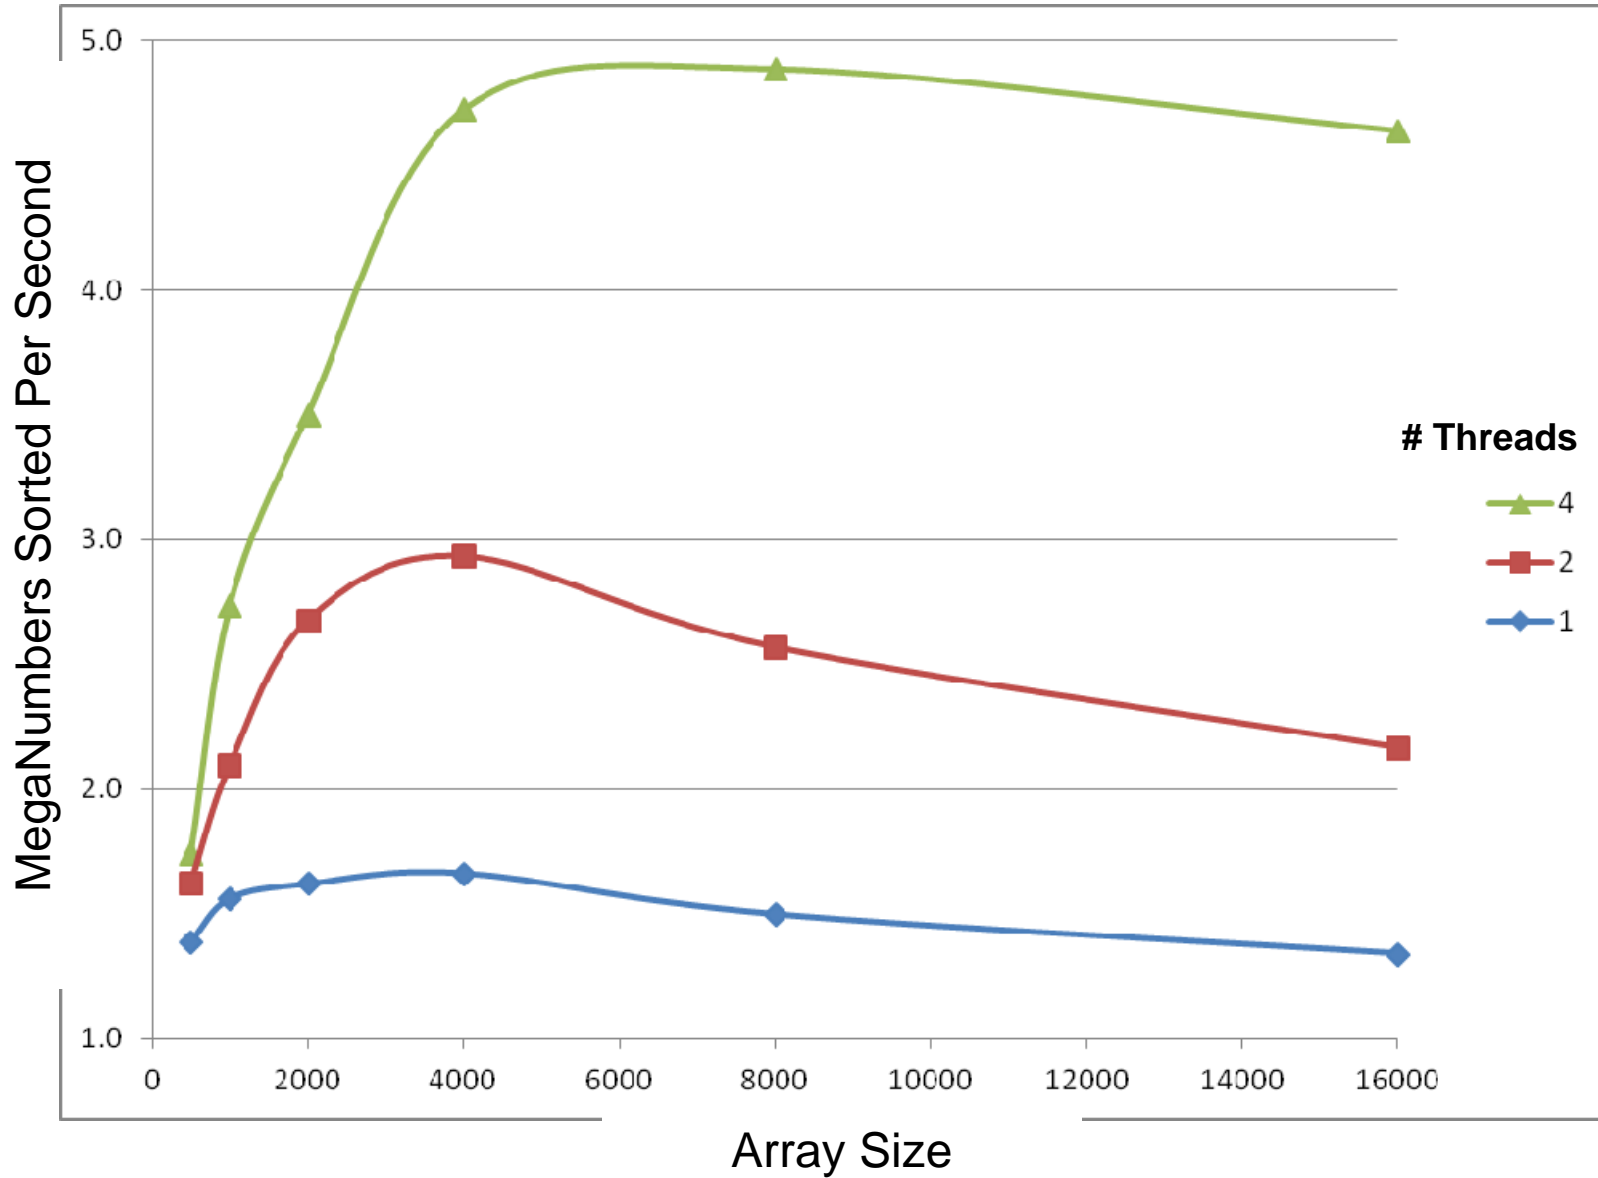
**Non-threaded**

original	Step #					
	0	1	2	3	4	5
6	5	5	3	3	1	1
5	6	3	5	1	3	2
4	3	6	1	5	2	3
3	4	1	6	2	5	4
2	1	4	2	6	4	5
1	2	2	4	4	6	6

**Threaded**

# OpenMP Performance as a Function of Array Size

7



# OpenMP Performance as a Function of # of Threads

8

