## **OpenMP Case Study: Trapezoid Integration Example**

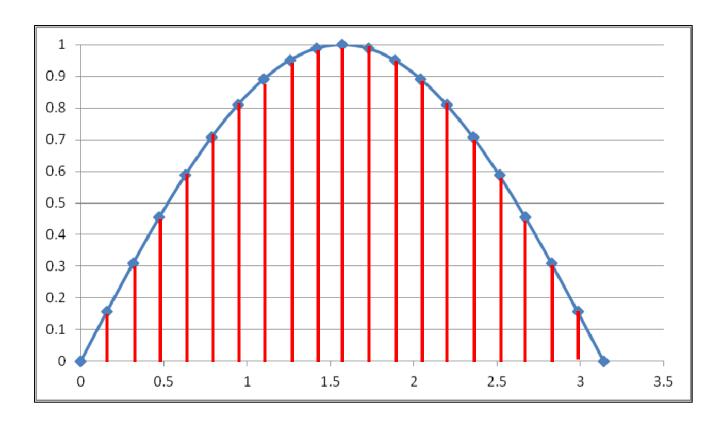
### **Mike Bailey**

mjb@cs.oregonstate.edu

### **Oregon State University**



# Find the area under the curve $y = \sin(x)$ for $0 \le x \le \pi$ using the Trapezoid Rule



Exact answer:  $\int_0^{\pi} (\sin x) dx = -\cos x \Big|_0^{\pi} = 2.0$ 



### Don't do it this way!

```
const double A = 0.;
const double B = M PI;
double dx = (B - A) / (float) (numSubdivisions - 1);
double sum = ( Function(A) + Function(B) ) / 2.;
omp_set_num_threads( numThreads );
#pragma omp parallel for default(none), shared(dx, sum)
for( int i = 1; i < numSubdivisions - 1; i++)
         double x = A + dx * (float) i;
         double f = Function(x);
         sum += f;
sum *= dx;
```

#### **Assembly code:**

Load sum
Add f
Store sum

What if the scheduler decides to switch threads right here?

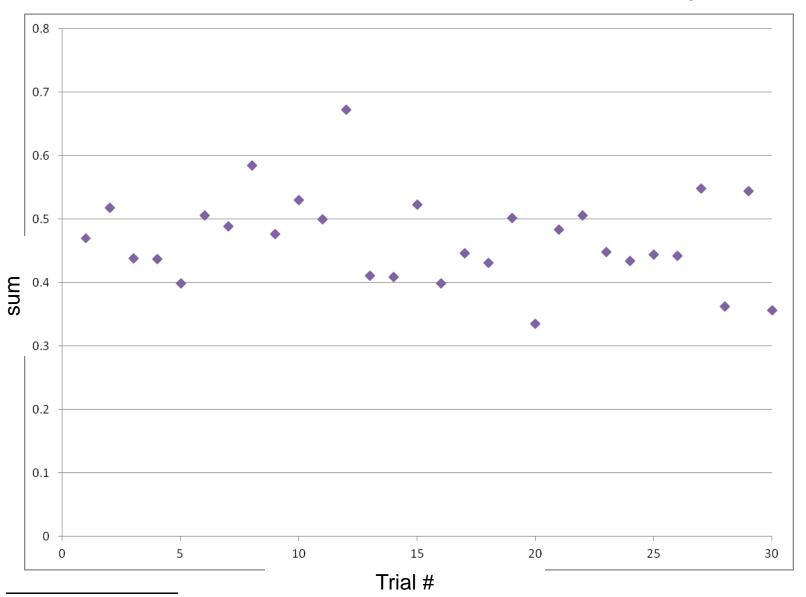


### The answer should be 2.0 *exactly*, but in 30 trials, it's not even close.<sup>4</sup> And, the answers aren't even consistent. Why?

0.469635	0.398893
0.517984	0.446419
0.438868	0.431204
0.437553	0.501783
0.398761	0.334996
0.506564	0.484124
0.489211	0.506362
0.584810	0.448226
0.476670	0.434737
0.530668	0.444919
0.500062	0.442432
0.672593	0.548837
0.411158	0.363092
0.408718	0.544778
0.523448	0.356299



### The answer should be 2.0 *exactly*, but in 30 trials, it's not even close.<sup>5</sup> And, the answers aren't even consistent. Why?



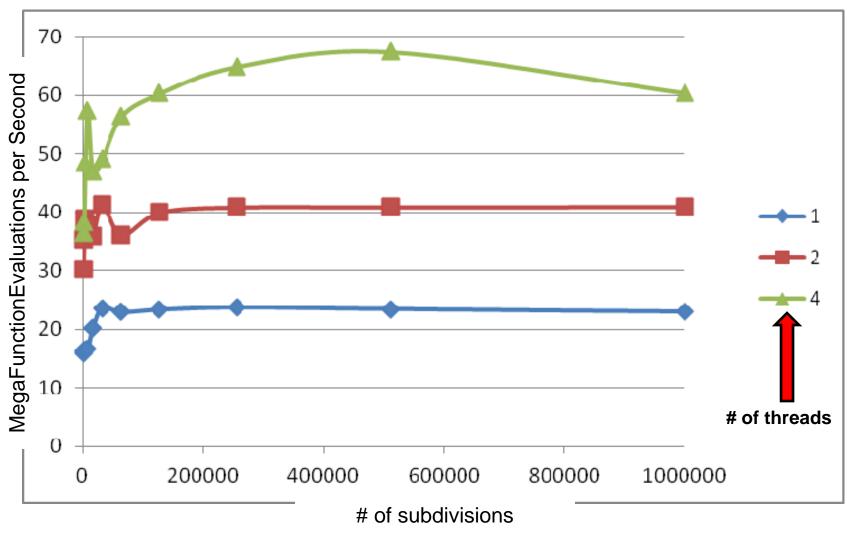


Oregon State University Computer Graphics

### Do it this way!

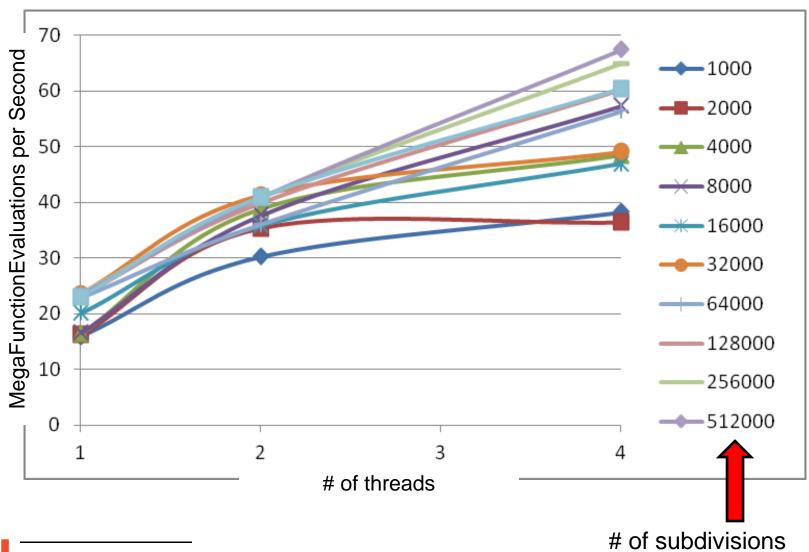
```
const double A = 0.;
const double B = M_PI;
double dx = (B - A) / (float) (numSubdivisions - 1);
omp_set_num_threads( numThreads );
double sum = ( Function(A) + Function(B) ) / 2.;
#pragma omp parallel for default(none), shared(dx), reduction(+:sum)
for( int i = 1; i < numSubdivisions - 1; i++)
         double x = A + dx * (float) i;
         double f = Function(x);
         sum += f;
sum *= dx;
```

### MegaFunctionEvaluations Per Second vs. Number of Subdivisions <sup>7</sup>





### MegaFunctionEvaluations Per Second vs. Number of Threads



### Ways to Make the Summing Work: Reduction vs. Atomic vs. Critical

```
#pragma omp parallel for shared(dx),reduction(+:sum)
for( int i = 0; i < numSubdivisions; i++ )
{
    double x = A + dx * (float) i;
    double f = Function(x);
    sum += f;
}</pre>
```

```
#pragma omp parallel for shared(dx)
for( int i = 0; i < numSubdivisions; i++ )
{
    double x = A + dx * (float) i;
    double f = Function(x);

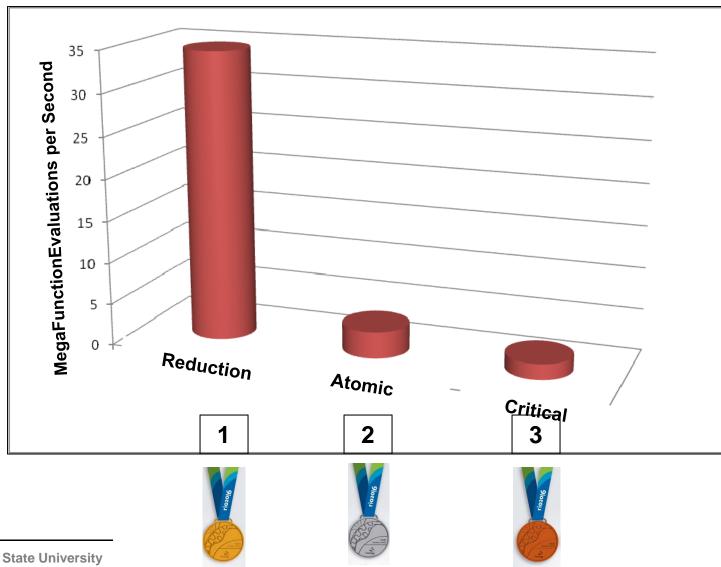
    #pragma omp atomic
    sum += f;
}</pre>
2
```

```
#pragma omp parallel for shared(dx)
for( int i = 0; i < numSubdivisions; i++ )
{
    double x = A + dx * (float) i;
    double f = Function(x);
    #pragma omp critical
    sum += f;
}</pre>
```

3

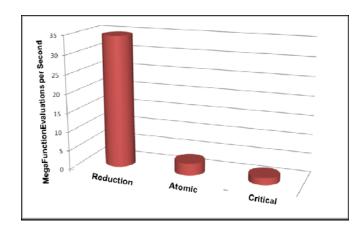


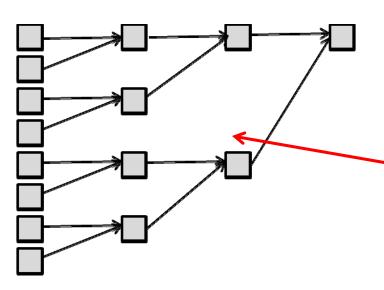
### Speed of using Reduction vs. Atomic vs. Critical



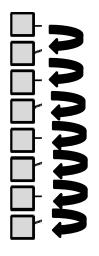


Oregon State University Computer Graphics

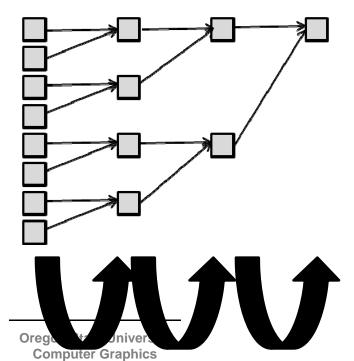




- Reduction secretly creates a temporary private variable for each thread's running sum. Each thread adding into its running sum doesn't interfere with any other thread adding into its running sum, and so threads don't need to slow down to get out of the way of each other.
- Reduction automatically creates a binary tree structure, like this, to add the N running sums in log<sub>2</sub>N time instead N time.



Serial addition: 8 numbers requires 7 steps



Parallel addition: 8 numbers requires 3 steps