



Tree Traversal using OpenMP Tasks

Mike Bailey

mjb@cs.oregonstate.edu


Oregon State University



Tree Traversal Algorithms

2

Given a tree:



**We would like to traverse it as quickly as possible.
We are assuming that we do not need to traverse it in order.
We just need to visit all nodes.**

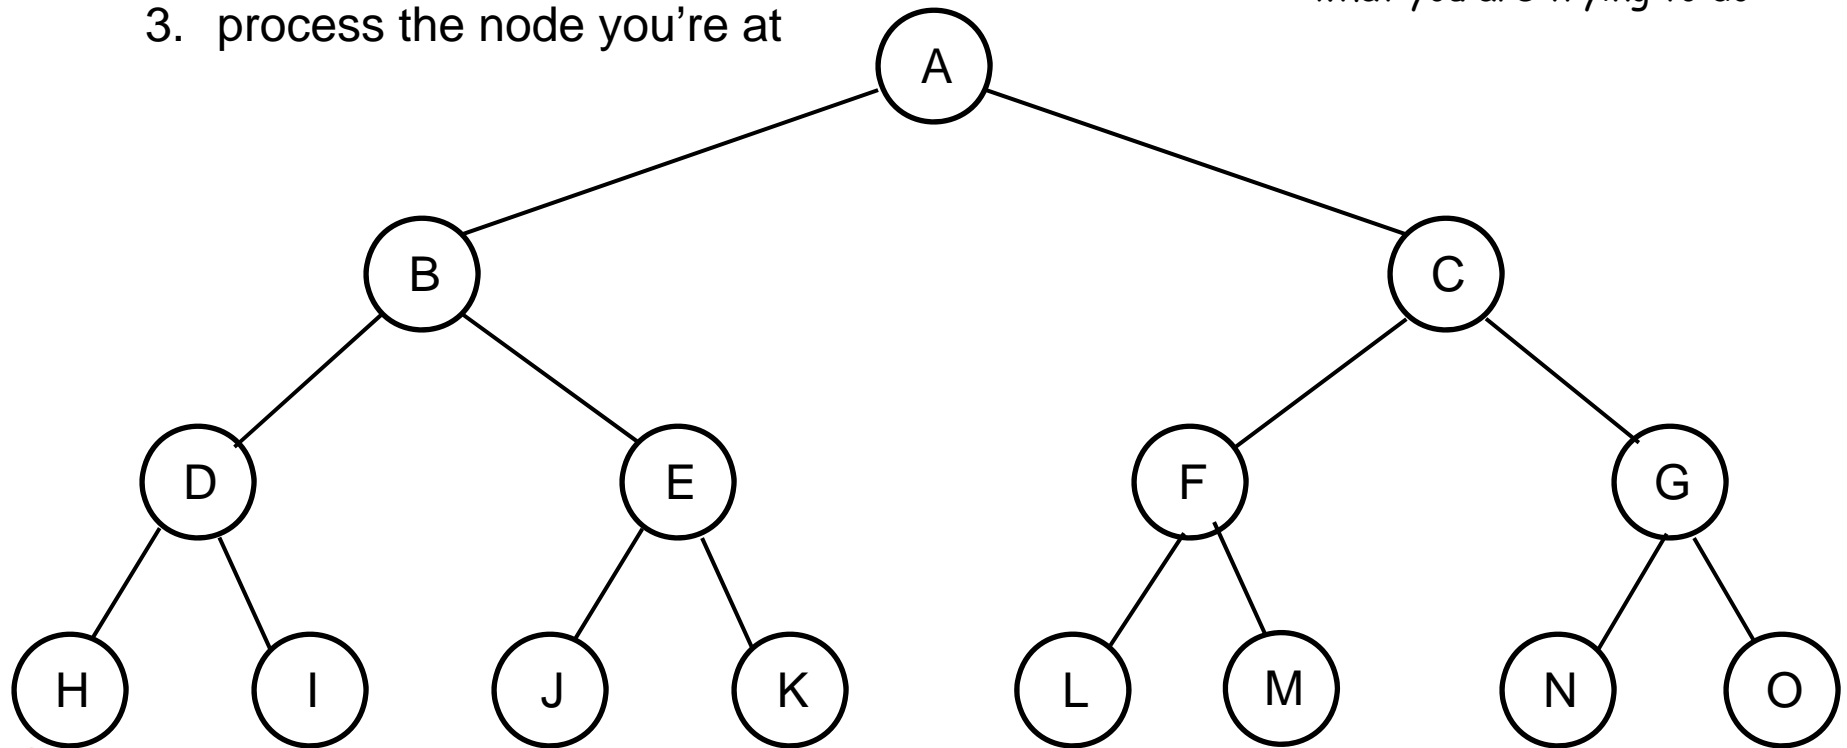
Tree Traversal Algorithms

- Common in graph algorithms, such as searching.
- If the tree is binary and is balanced, then the maximum depth of the tree is $\log_2(\# \text{ of Nodes})$

- Strategy at a node:

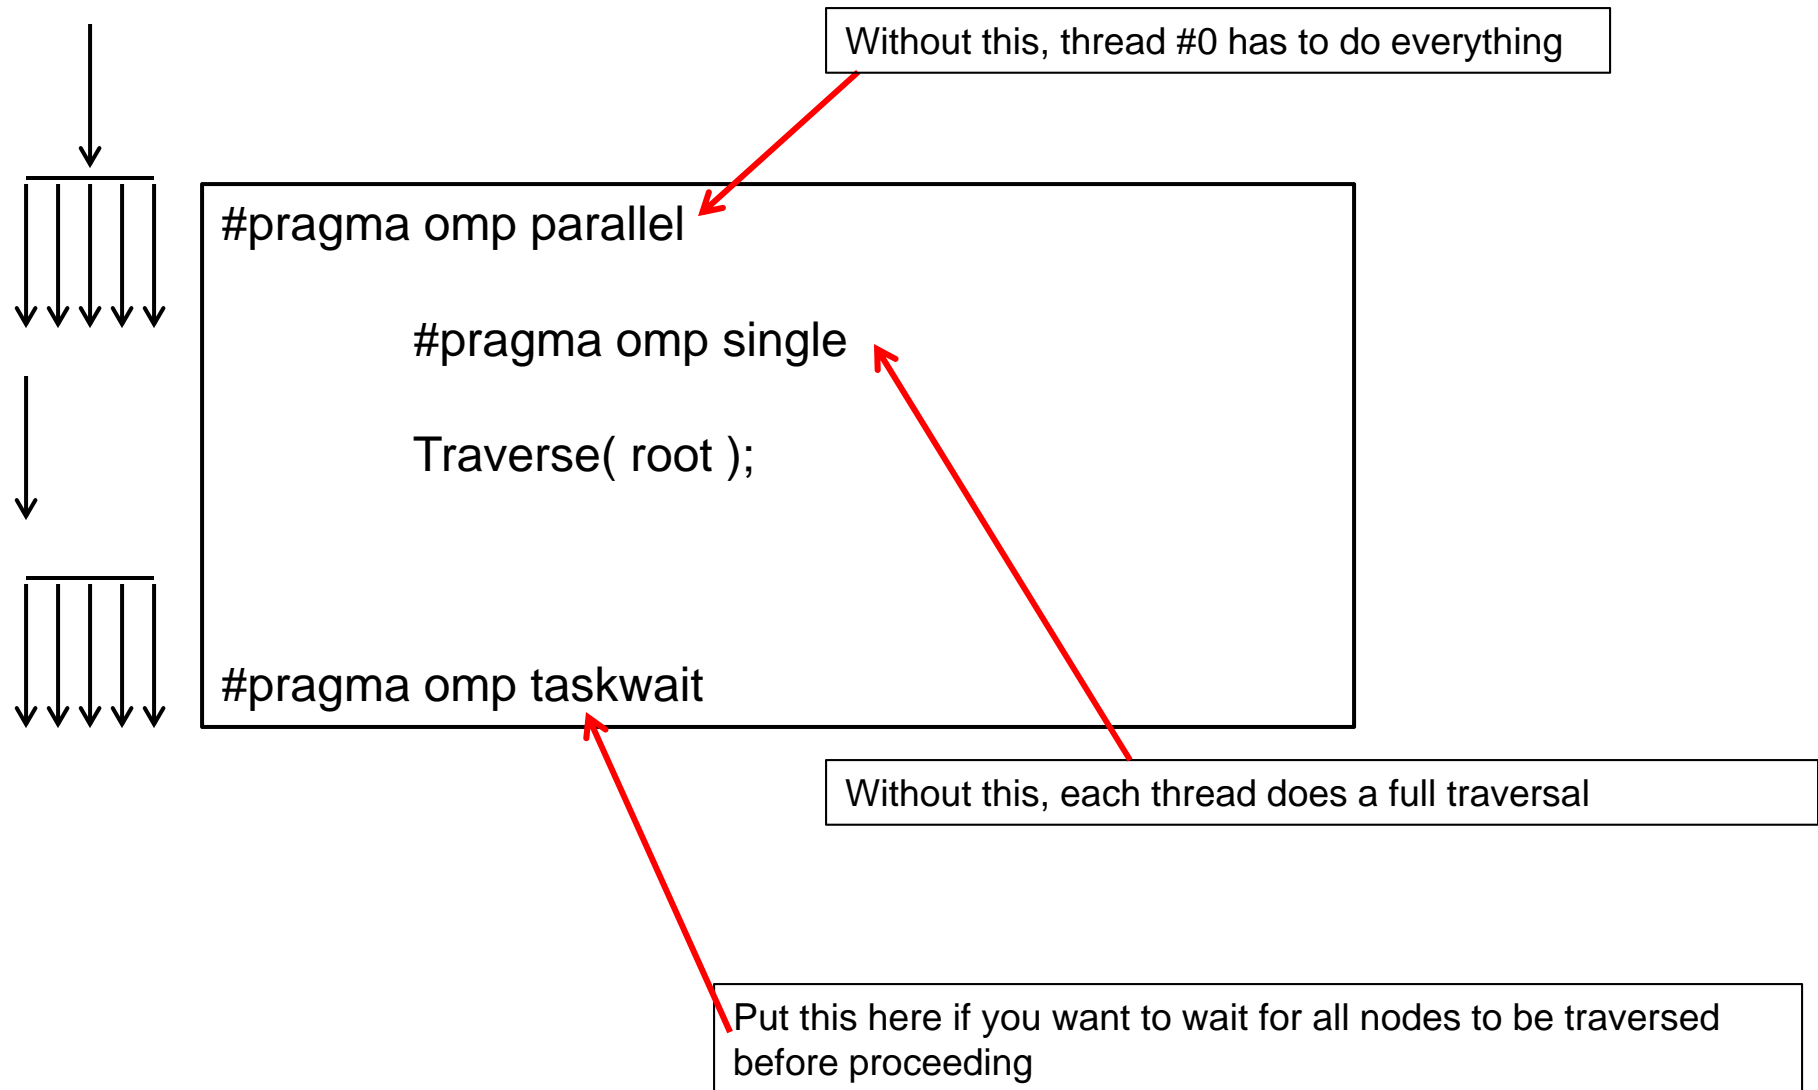
1. follow one descendent node
2. follow the other descendent node
3. process the node you're at

This order could be re-arranged, depending on what you are trying to do





Tree Traversal Algorithms

4



Parallelizing a Binary Tree Traversal with Tasks

5



```
void
Traverse( Node *n )
{
    if( n->left != NULL )
    {
        #pragma omp task private(n) untied
        Traverse( n->left );
    }

    if( n->right != NULL )
    {
        #pragma omp task private(n) untied
        Traverse( n->right );
    }

    #pragma omp taskwait

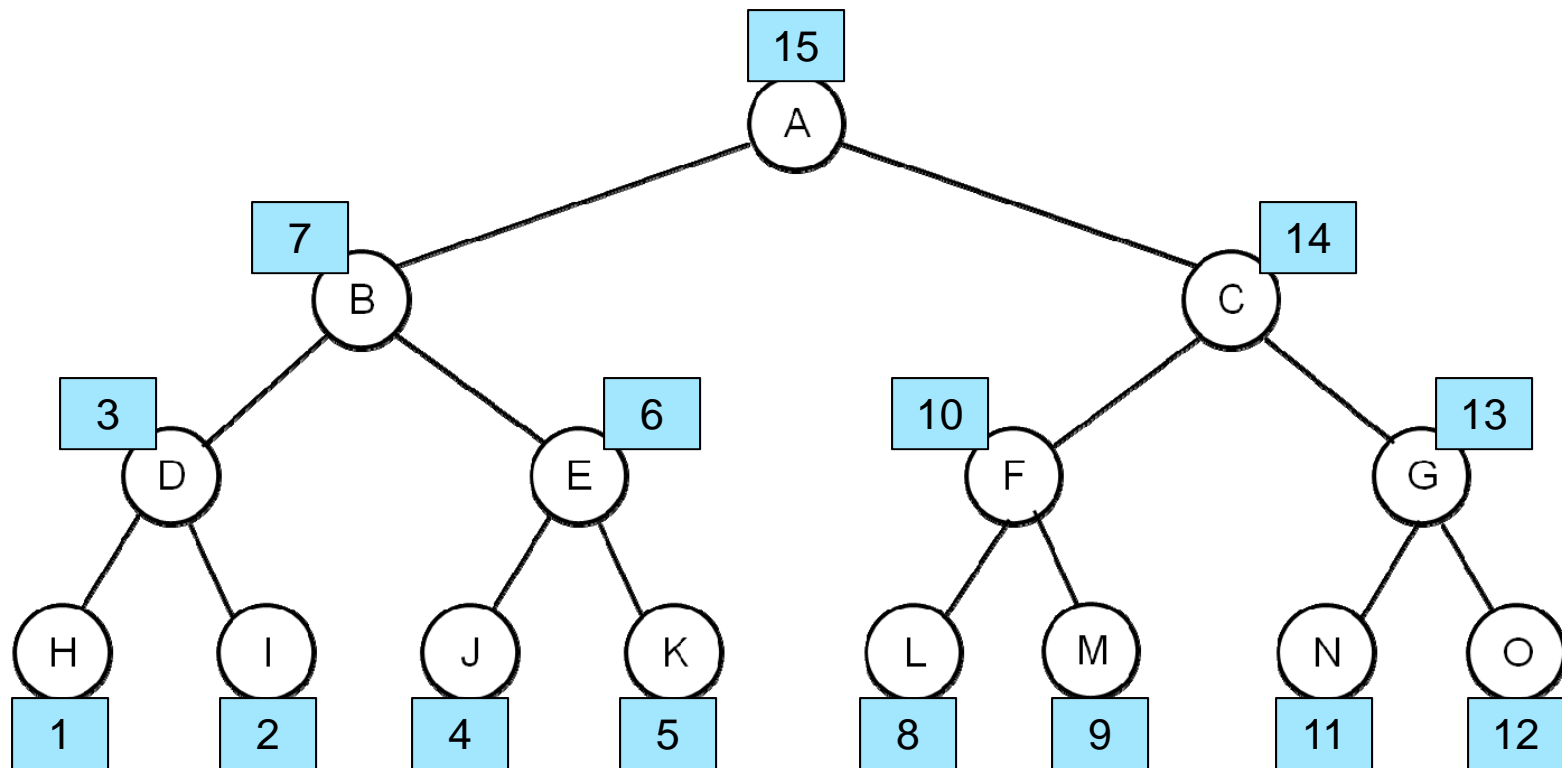
    Process( n );
}
```

Put this here if you want to wait for both branches to be taken before processing the parent

Parallelizing a Binary Tree Traversal with Tasks

6

Traverse(A);

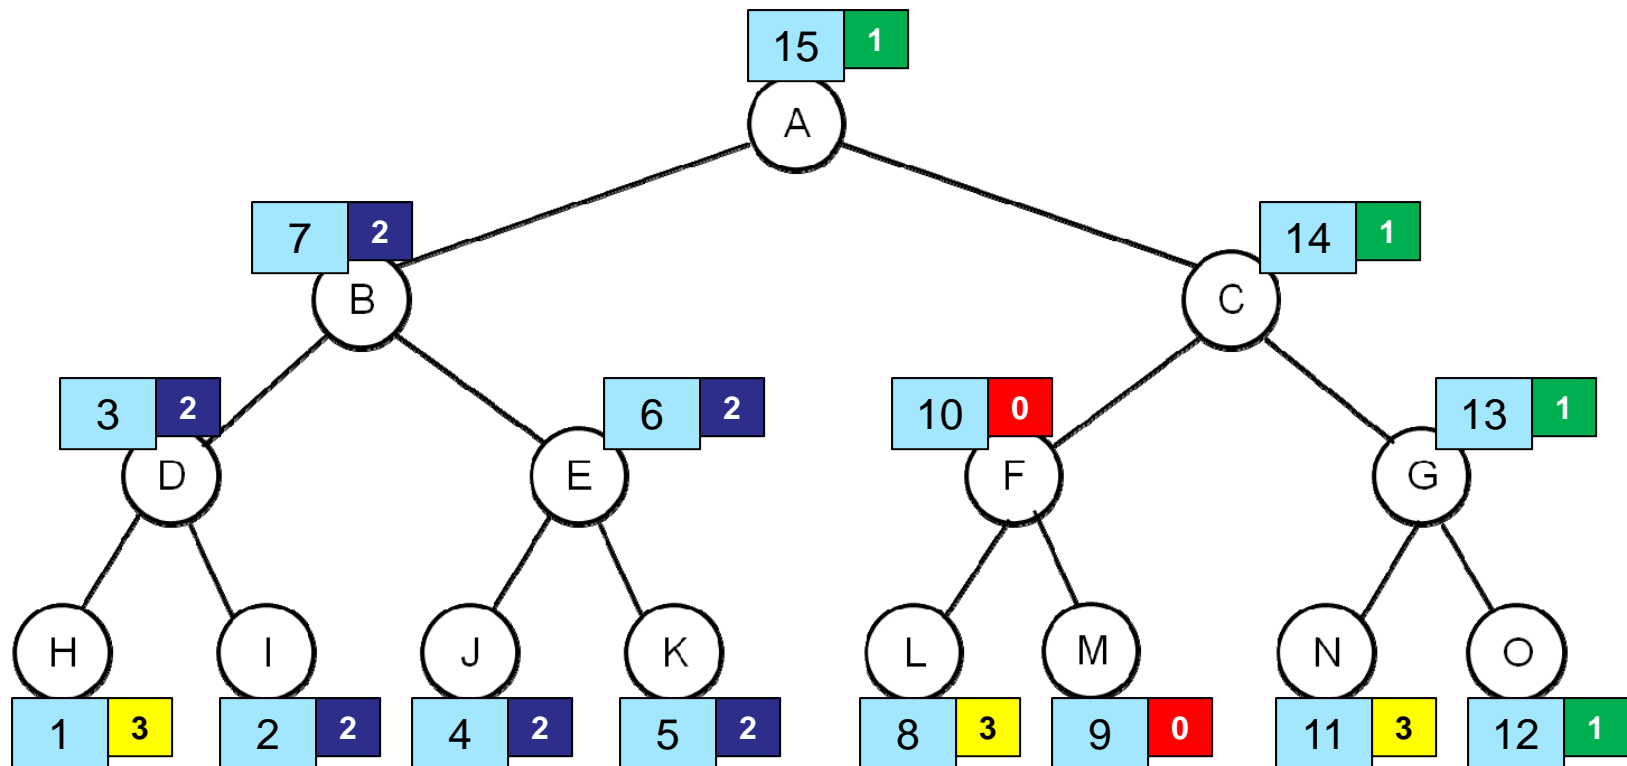


Parallelizing a Binary Tree Traversal with Tasks: *Tied*

7



Traverse(A);

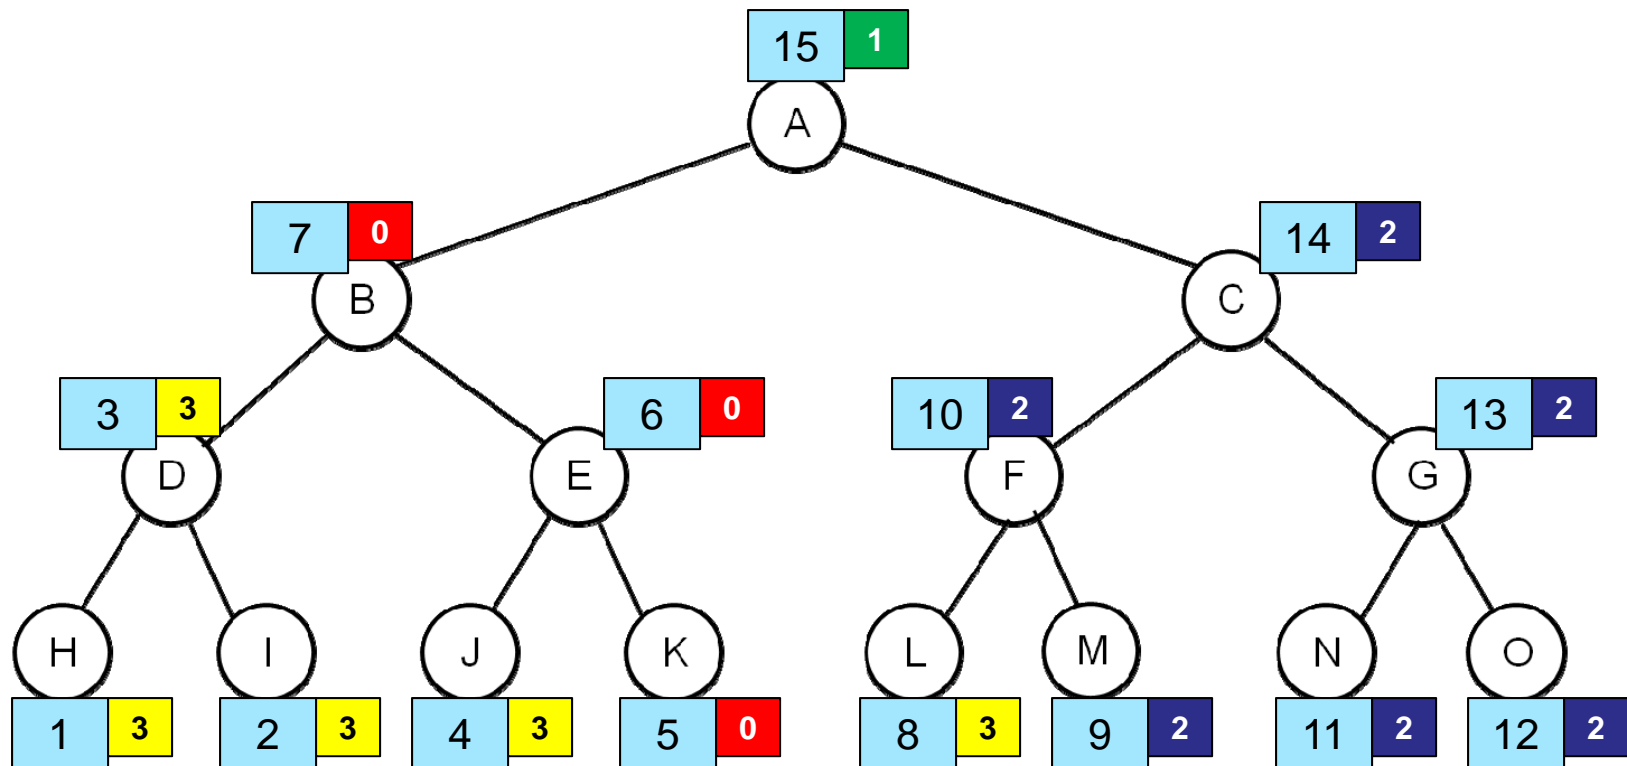


Parallelizing a Binary Tree Traversal with Tasks: *Untied*

8



Traverse(A);



How Evenly Tasks Get Assigned to Threads

9

6 Levels – g++ 4.9:

Thread #	Number of Tasks
0	1
1	32
2	47
3	47

6 Levels – icpc 15.0.0:

Thread #	Number of Tasks
0	29
1	31
2	41
3	26

12 Levels – g++ 4.9:

Thread #	Number of Tasks
0	2561
1	2
2	2813
3	2815

12 Levels – icpc 15.0.0:


Thread #	Number of Tasks
0	1999
1	2068
2	2035
3	2089

Benchmarking a Binary Task-driven Tree Traversal


10

```
void
Process( Node *n )
{
    for( int i = 0; i < 1024; i++ )
    {
        n->value = pow( n->value, 1.1 );
    }
}
```

Performance vs. Number of Threads



Performance vs. Number of Levels



- Tasks get spread among the current “thread team”
- Tasks can execute immediately or can be deferred. They are executed at “some time”.
- Tasks can be moved between threads, that is, if one thread has a backlog of tasks to do, an idle thread can come steal some workload.
- Tasks are more dynamic than sections. The task paradigm would still work if there was a variable number of children at each node.

Parallelizing an N-Tree Traversal with Tasks


14

```
void
Traverse( Node *n )
{
    for( int i = 0; i < n->numChildren; i++ )
    {
        if( n->child[i] != NULL )
        {
            #pragma omp task
            Traverse( n->child[i] );
        }
    }

    #pragma omp taskwait

    Process( n );
}
```

Performance vs. Number of Levels




8-thread Speed-up ≈ 6.7

$F_p \approx ??\%$

Max Speed-up $\approx ??$

Performance vs. Number of Threads



8-thread Speed-up ≈ 6.7

$F_p \approx ??\%$

Max Speed-up $\approx ??$

$$F_p = \frac{n}{(n-1)} \left(1 - \frac{1}{Speedup} \right) = 97\%$$

$$\max Speedup = \frac{1}{1 - F_p} = 33x$$