


Vector Processing

(aka, Single Instruction Multiple Data, or SIMD)

Mike Bailey

mjb@cs.oregonstate.edu

Oregon State University



What is Vectorization/SIMD and Why do We Care?


2

Performance!

Many hardware architectures today, both CPU and GPU, allow you to perform arithmetic operations on multiple array elements simultaneously.

(Thus the label, “Single Instruction Multiple Data”.)

We care about this because many problems, especially scientific and engineering, can be cast this way. Examples include convolution, Fourier transform, power spectrum, autocorrelation, etc.



SIMD in Intel Chips

3

Year Released	Name	Width (bits)	Width (FP words)
1996	MMX	64	2
1999	SSE	128	4
2011	AVX	256	8
2013	AVX-512	512	16

Xeon Phi

Note: one complete cache line!

Intel SSE

4


Year Released	Name	Width (bits)	Width (FP words)
1996	MMX	64	2
1999	SSE	128	4
2011	AVX	256	8
2013	AVX-512	512	16

Intel architecture supports vectorization. The most well-known form is called Streaming SIMD Extension, or **SSE**. It allows four floating point operations to happen simultaneously.

Normally a *scalar* floating point multiplication instruction happens like this:

mulss r1, r0


← “ATT form”:
mulss src, dst



The SSE version of the multiplication instruction happens like this:

mulps xmm1, xmm0

← “ATT form”:
mulps src, dst



Array * Array

```
void  
SimdMul( float *a, float *b, float *c, int len )  
{  
    c[0:len] = a[0:len] * b[0:len];  
}
```

Note that the construct:

a[0 : ArraySize]

Is meant to be read as:

“The set of elements in the array **a** starting at index 0 and going for **ArraySize** elements”.

not as:

“The set of elements in the array **a** starting at index 0 and going through index **ArraySize**”.

Array * Array

```
void  
SimdMul( float *a, float *b, float *c, int len )  
{  
    #pragma omp simd  
    for( int i= 0; i < len; i++ )  
        c[i] = a[i] * b[i];  
}
```

Array * Scalar


```
void  
SimdMul( float *a, float b, float *c, int len )  
{  
    c[0:len] = a[0:len] * b;  
}
```

Array * Scalar

```
void  
SimdMul( float *a, float b, float *c, int len )  
{  
    #pragma omp simd  
    for( int i = 0; i < len; i++ )  
        c[i] = a[i] * b;  
}
```


Array*Array Multiplication Speed

8



Array*Array Multiplication Speedup

9



You would think it would always be $4.0 \pm$ noise effects, but it's not. Why?

Combining SIMD with Parallel Threading in OpenMP 4.0

10

```
#pragma omp parallel for
```

```
for( int i = 0; i < ArraySize; i++ )  
{  
    c[ i ] = a[ i ] * b[ i ];  
}
```

```
#pragma omp simd
```

```
for( int i = 0; i < ArraySize; i++ )  
{  
    c[ i ] = a[ i ] * b[ i ];  
}
```

```
#pragma omp parallel for simd
```

```
for( int i = 0; i < ArraySize; i++ )  
{  
    c[ i ] = a[ i ] * b[ i ];  
}
```

Requirements for a For-Loop to be Vectorized

11

- If there are nested loops, the one to vectorize must be the inner one.
- There can be no jumps or branches. “Masked assignments” (an if-statement-controlled assignment) are OK, e.g.,

```
if( A[ i ] > 0. )  
    B[ i ] = 1.;
```

- The total number of iterations must be known at runtime when the loop starts
- There cannot be any backward loop dependencies, like this:
$$A[i] = A[i-1] + 1.;$$
- It helps if the elements have contiguous memory addresses.

Prefetching is used to place a cache line in memory before it is to be used, thus hiding the latency of fetching from off-chip memory.

There are two key issues here:

1. Issuing the prefetch at the right time
2. Issuing the prefetch at the right distance

The right time:

If the prefetch is issued too late, then the memory values won't be back when the program wants to use them, and the processor has to wait anyway.

If the prefetch is issued too early, then there is a chance that the prefetched values could be evicted from cache by another need before they can be used.

The right distance:

The “prefetch distance” is how far ahead the prefetch memory is than the memory we are using right now.

Too far, and the values sit in cache for too long, and possibly get evicted.

Too near, and the program is ready for the values before they have arrived.

Array Multiplication

Length of Arrays (NUM): 1,000,000


Length per SIMD call (ONETIME): 256

```
for( int i = 0; i < NUM; i += ONETIME )
{
    __builtin_prefetch ( &A[i+PD], WILL_READ_ONLY,          LOCALITY_LOW );
    __builtin_prefetch ( &B[i+PD], WILL_READ_ONLY,          LOCALITY_LOW );
    __builtin_prefetch ( &C[i+PD], WILL_READ_AND_WRITE, LOCALITY_LOW );

    SimdMul( A, B, C, ONETIME );
}
```

The Effects of Prefetching on SIMD Computations

14



This all sounds great! What is the catch?

The catch is that compilers haven't caught up to producing efficient SIMD code. So, while there are great ways to express the desire for SIMD in code, you won't get the full potential speedup ... yet.

So, for the CPU SIMD project, we are going to investigate the potential speedup using assembly language. Don't worry – *you* don't need to write it.

You will be given two assembly functions:

1. SimdMul: $C[0:len] = A[0:len] * B[0:len]$
2. SimdMulSum: $\text{return} (\sum A[0:len] * B[0:len])$

Warning – due to the nature of how different compilers and systems handle local variables, these two functions only work on *flip* using gcc/g++, without `-O3` !!!

Getting at the full SIMD power until compilers catch up

16

```
void
SimdMul( float *a, float *b, float *c, int len )
{
    int limit = ( len/SSE_WIDTH ) * SSE_WIDTH;
    __asm
    (
        ".att_syntax\n\t"
        "movq  -24(%rbp), %rbx\n\t"      // a
        "movq  -32(%rbp), %rcx\n\t"      // b
        "movq  -40(%rbp), %rdx\n\t"      // c
    );

    for( int i = 0; i < limit; i += SSE_WIDTH )
    {
        __asm
        (
            ".att_syntax\n\t"
            "movups (%rbx), %xmm0\n\t"    // load the first sse register
            "movups (%rcx), %xmm1\n\t"    // load the second sse register
            "mulps  %xmm1, %xmm0\n\t"      // do the multiply
            "movups %xmm0, (%rdx)\n\t"     // store the result
            "addq $16, %rbx\n\t"
            "addq $16, %rcx\n\t"
            "addq $16, %rdx\n\t"
        );
    }

    for( int i = limit; i < len; i++ )
    {
        c[i] = a[i] * b[i];
    }
}
```

**This only works on *flip* using gcc/g++,
without `-O3` !!!**

Getting at the full SIMD power until compilers catch up

17

```
float
SimdMulSum( float *a, float *b, int len )
{
    float sum[4] = { 0., 0., 0., 0. };
    int limit = ( len/SSE_WIDTH ) * SSE_WIDTH;

    __asm
    (
        ".att_syntax\n\t"
        "movq  -40(%rbp), %rbx\n\t"      // a
        "movq  -48(%rbp), %rcx\n\t"      // b
        "leaq  -32(%rbp), %rdx\n\t"      // &sum[0]
        "movups (%rdx), %xmm2\n\t"      // 4 copies of 0. in xmm2
    );

    for( int i = 0; i < limit; i += SSE_WIDTH )
    {
        __asm
        (
            ".att_syntax\n\t"
            "movups (%rbx), %xmm0\n\t"    // load the first sse register
            "movups (%rcx), %xmm1\n\t"    // load the second sse register
            "mulps  %xmm1, %xmm0\n\t"      // do the multiply
            "addps  %xmm0, %xmm2\n\t"      // do the add
            "addq $16, %rbx\n\t"
            "addq $16, %rcx\n\t"
        );
    }

    __asm
    (
        ".att_syntax\n\t"
        "movups  %xmm2, (%rdx)\n\t"      // copy the sums back to sum[ ]
    );

    for( int i = limit; i < len; i++ )
    {
        sum[i-limit] += a[i] * b[i];
    }

    return sum[0] + sum[1] + sum[2] + sum[3];
}
```

**This only works on *flip* using gcc/g++,
without -O3 !!!**

A preview of things to come: OpenCL has a data type called “float4”

When we get to OpenCL, we could compute projectile physics like this:

```
float4  pp;                                // p'  
pp.x    = p.x + v.x*DT;  
pp.y    = p.y + v.y*DT + .5*DT*DT*G.y;  
pp.z    = p.z + v.z*DT;
```

But, instead, we will do it like this:

```
float4  pp = p + v*DT + .5*DT*DT*G;        // p'
```

We do it this way for two reasons:

1. Convenience and clean coding
2. Some hardware can do multiple arithmetic operations simultaneously

A preview of things to come: OpenCL has a data type called “float4”

The whole thing will look like this:

```
constant float4 G      = (float4) ( 0., -9.8, 0., 0. );
constant float  DT      = 0.1;

kernel
void
Particle( global float4 * dPobj, global float4 * dVel, global float4 * dCobj )
{
    int gid  = get_global_id( 0 );           // particle #
    float4 p  = dPobj[gid];                  // particle #gid's position
    float4 v  = dVel[gid];                   // particle #gid's velocity

    float4  pp  = p + v*DT + .5*DT*DT*G;      // p'
    float4  vp  = v + G*DT;                  // v'

    dPobj[gid] = pp;
    dVel[gid]  = vp;
}
```