

The Open Computing Language (OpenCL)

Also go look at the files **first.cpp** and **first.cl** !

Mike Bailey

mjb@cs.oregonstate.edu

Oregon State University



- OpenCL consists of two parts: a C/C++-callable API and a C-ish programming language.
- The programming language can run on NVIDIA GPUs, AMD GPUs, Intel CPUs, Intel GPUs, mobile devices, and (supposedly) FPGAs (Field-Programmable Gate Arrays). But, OpenCL is at its best on compute devices with large amounts of **data parallelism**, which usually implies GPUs.
- You break your computational problem up into lots and lots of small pieces. Each piece gets farmed out to threads on the GPU.
- Each thread wakes up and is able to ask questions about where it lives in the entire collection of (thousands of) threads. From that, it can tell what it is supposed to be working on.
- OpenCL can share data, and interoperate with, OpenGL
- There is a JavaScript implementation of OpenCL, called WebCL
- There is a JavaScript implementation of OpenGL, called WebGL
- WebCL can share data, and interoperate with, WebGL
- The GPU does not have a stack, and so the OpenCL C-ish programming language cannot do recursion and cannot make function calls. It also can't use pointers.


Who Is Behind OpenCL?

Members of Khronos's OpenCL Working Group

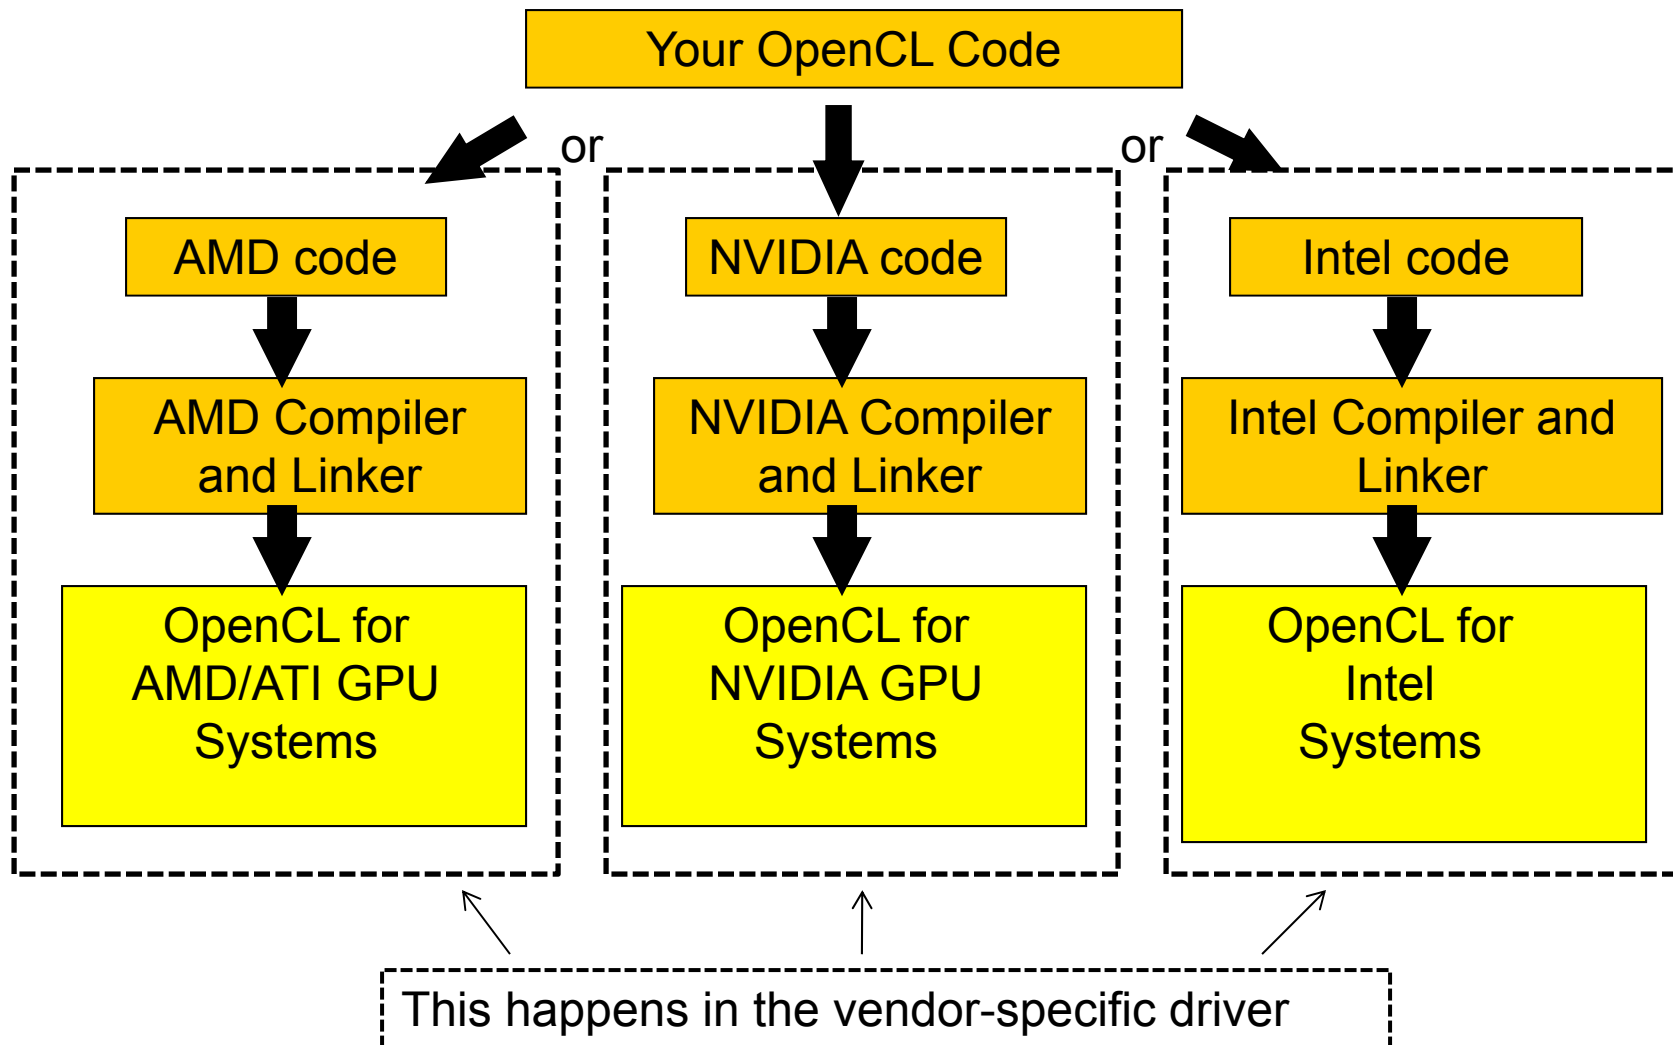


Example of using OpenCL in a System-on-a-Chip: Qualcomm Node – Full Linux and OpenCL

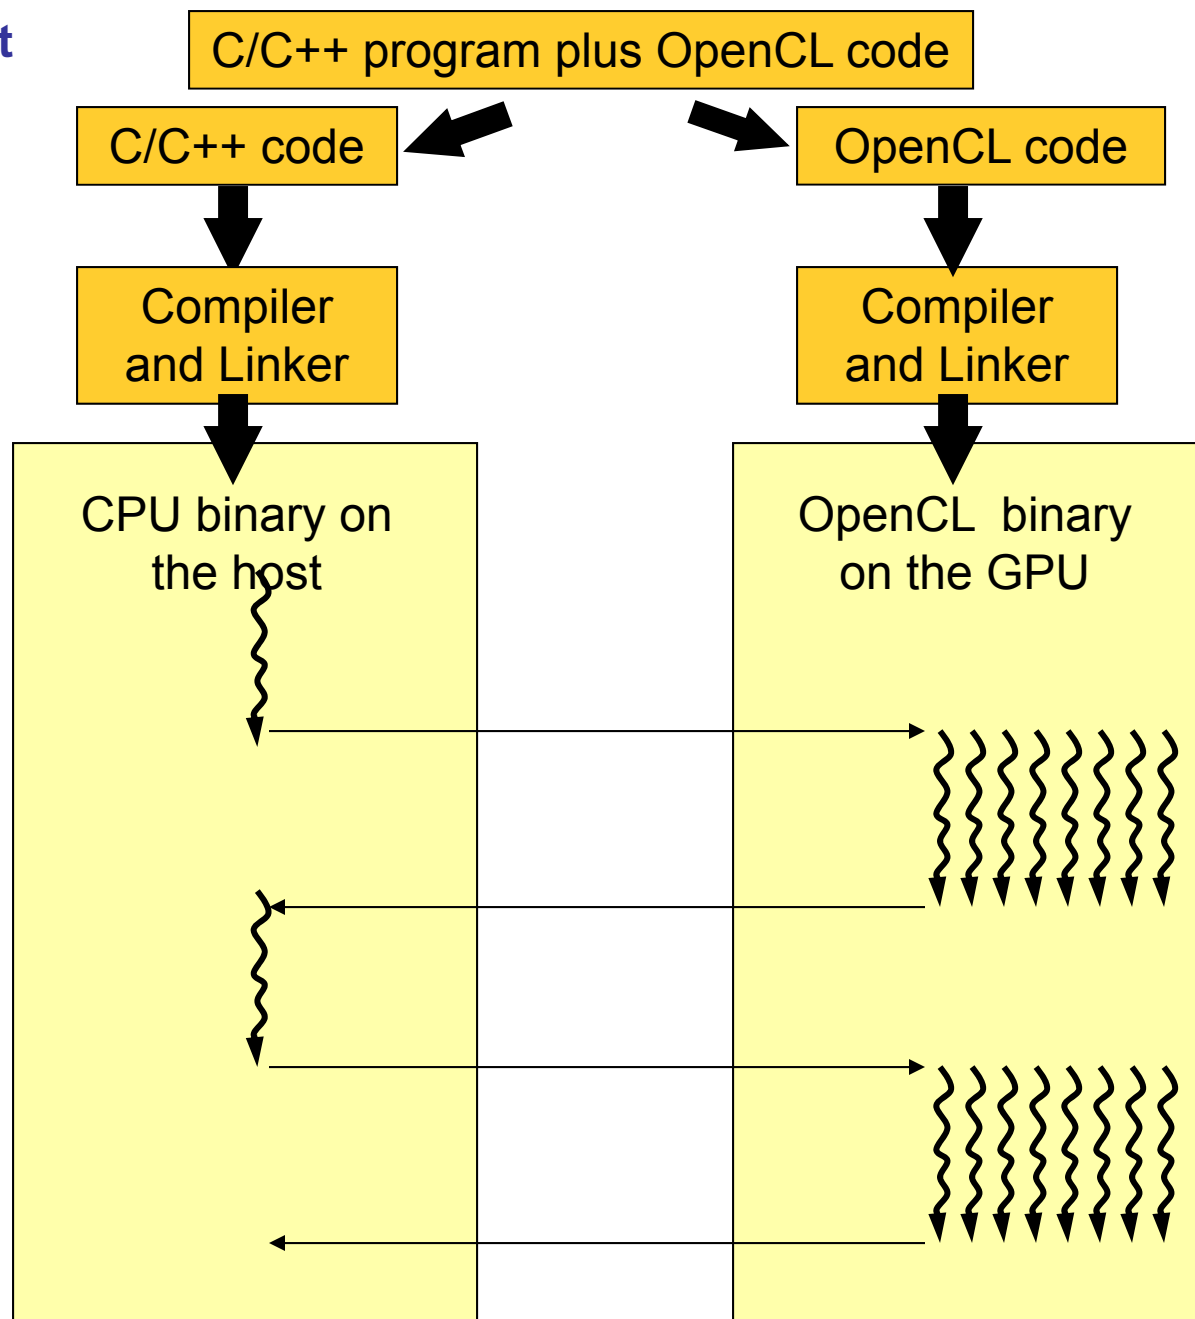
4



OpenCL – Vendor-independent GPU Programming



The OpenCL Programming Environment



OpenCL wants you to break the problem up into Pieces

If you were
writing in C/C++,
you would say:

```
void  
ArrayMult( int n, float *a, float *b, float *c)  
{  
    for ( int i = 0; i < n; i++ )  
        c[ i ] = a[ i ] * b[ i ];  
}
```

If you were writing
in OpenCL, you
would say:

```
kernel  
void  
ArrayMult( global float *dA, global float *dB, global float *dC)  
{  
    int  gid = get_global_id ( 0 );  
    dC[gid] = dA[gid] * dB[gid];  
}
```

The OpenCL Language also supports Vector Parallelism

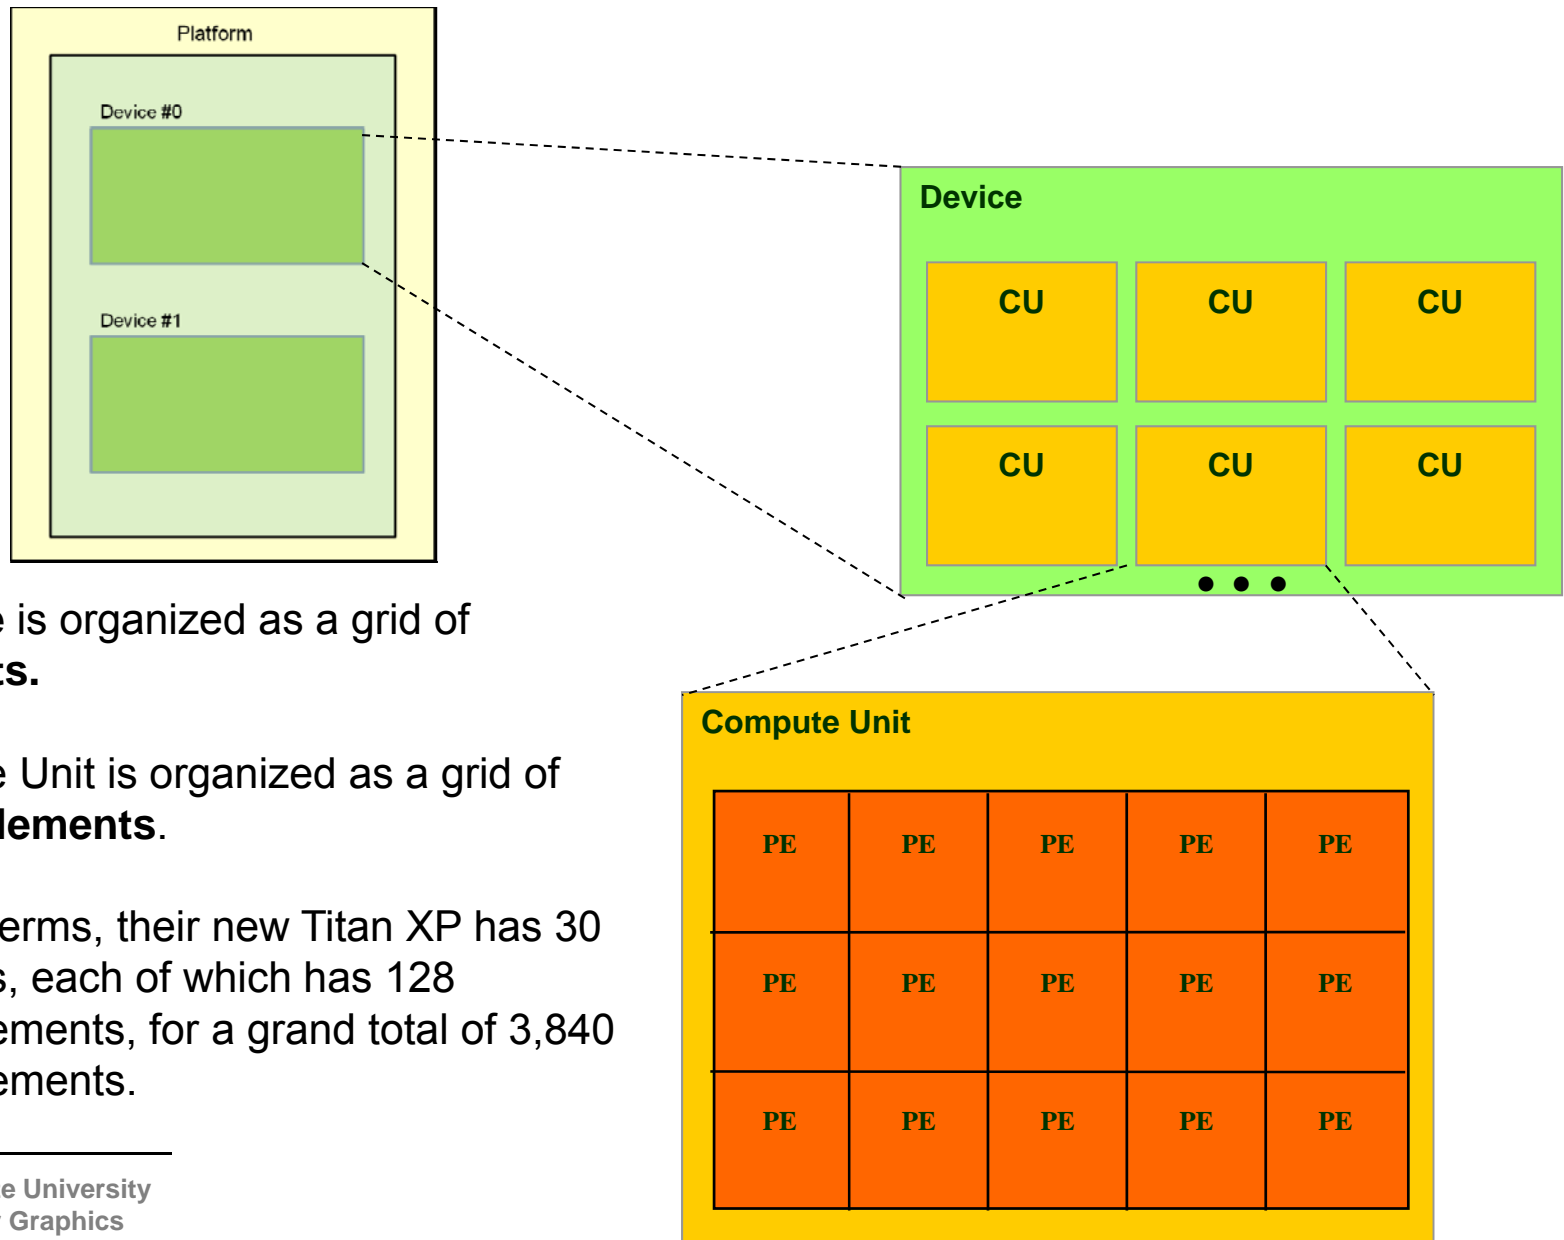
OpenCL code can be vector-oriented, meaning that it can perform a single instruction on multiple data values at the same time (SIMD).

Vector data types are: char_n , int_n , float_n , where $n = 2, 4, 8, \text{ or } 16$.

```
float4 f, g;  
f = (float4)( 1.f, 2.f, 3.f, 4.f );  
  
float16 a16, x16, y16, z16;  
  
f.x = 0.;  
f.xy = g.zw;  
x16.s89ab = f;  
  
float16 a16 = x16 * y16 + z16;
```

(Note: just because the language supports it, doesn't mean the hardware does.)

From the GPU101 Notes: Compute Units and Processing Elements are Arranged in Grids



A GPU **Device** is organized as a grid of **Compute Units**.

Each Compute Unit is organized as a grid of **Processing Elements**.

So in NVIDIA terms, their new Titan XP has 30 Compute Units, each of which has 128 Processing Elements, for a grand total of 3,840 Processing Elements.

OpenCL Software Terminology: Work-Groups and Work-Items are Arranged in Grids

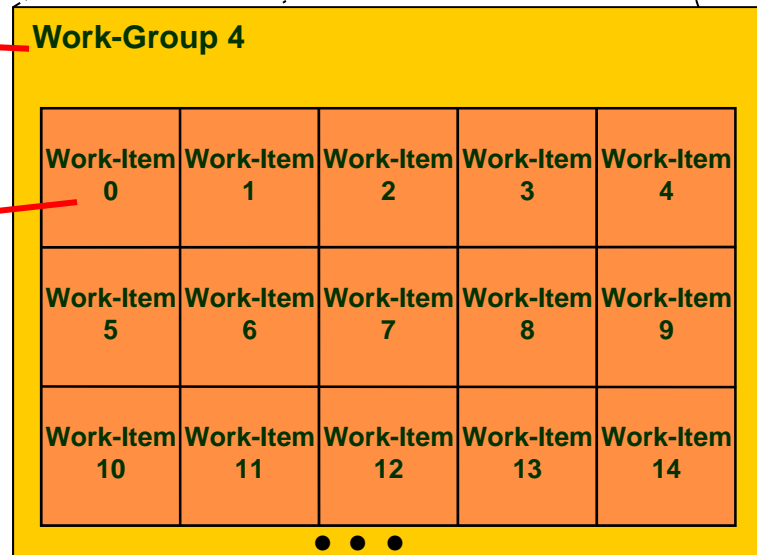


An OpenCL program is organized as a grid of **Work-Groups**.

Each Work-Group is organized as a grid of **Work-Items**.

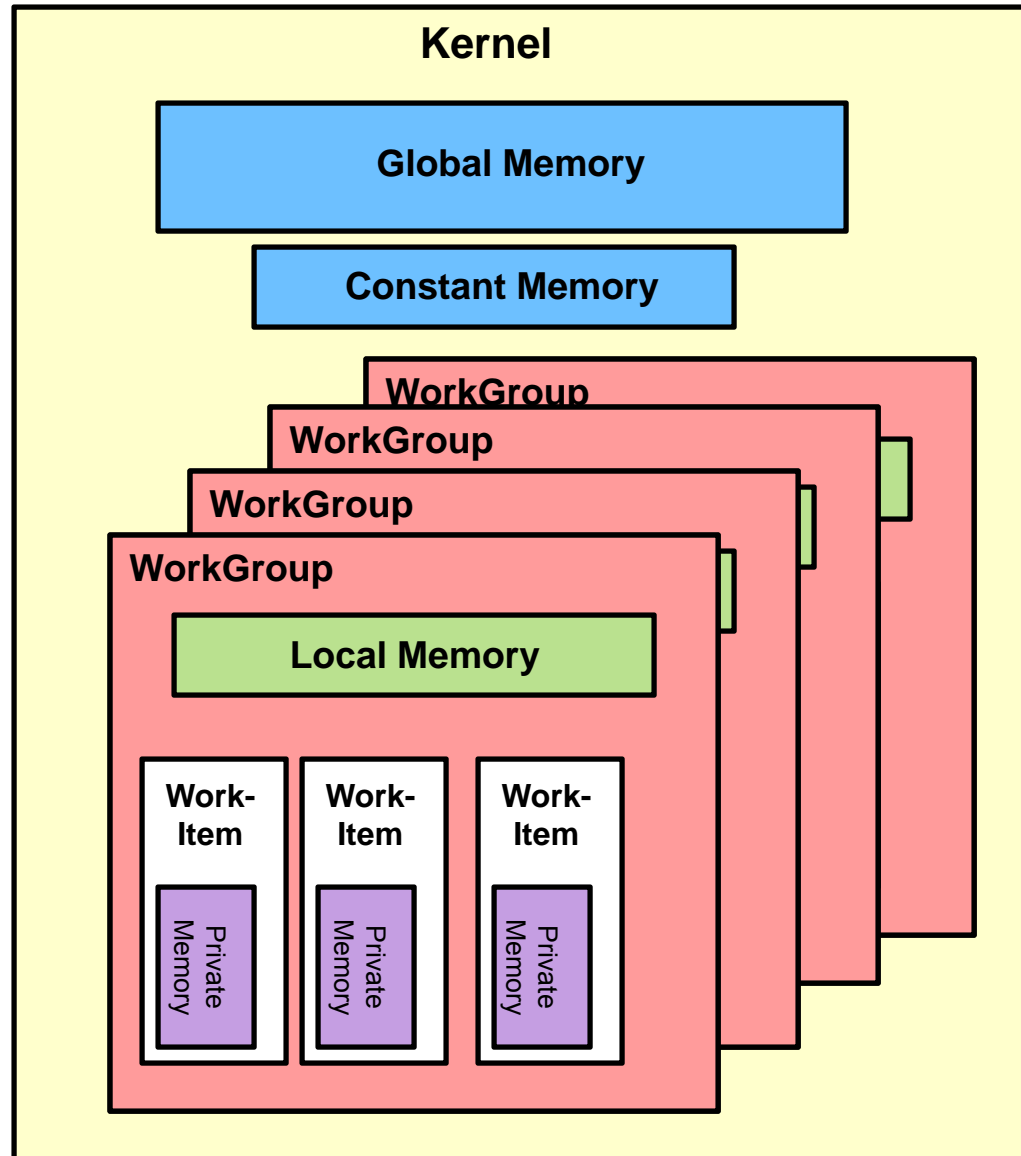
In terms of hardware, a Work-Group runs on a Compute Unit and a Work-Item runs on a Processing Element (PE).

One thread is assigned to each Work-Item.

Threads are swapped on and off the PEs.

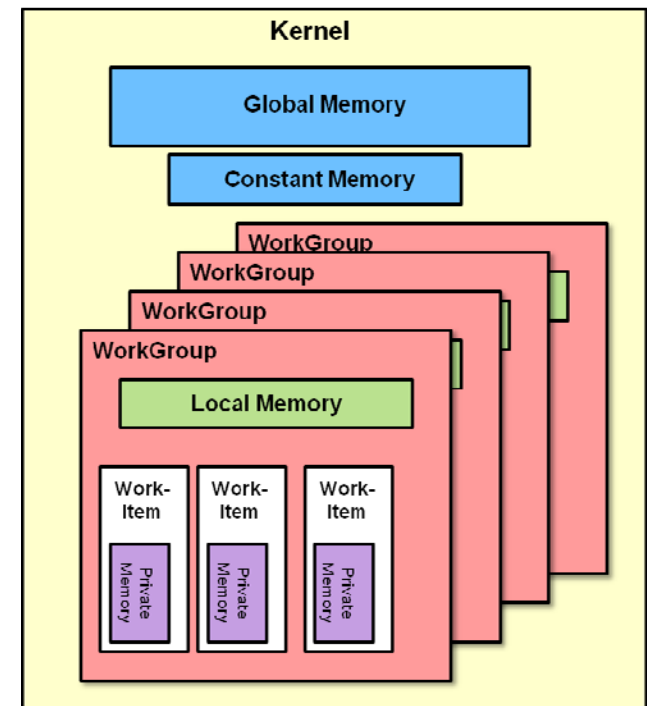


OpenCL Memory Model



Rules

- Threads can share memory with the other Threads in the same Work-Group
- Threads can synchronize with other Threads in the same Work-Group
- Global and Constant memory is accessible by all Threads in all Work-Groups
- Global and Constant memory is often cached inside a Work-Group
- Each Thread has registers and private memory
- Each Work-Group has a maximum number of registers it can use. These are divided equally among all its Threads



Querying the Number of Platforms (usually one)


```

cl_uint numPlatforms;
status = clGetPlatformIDs( 0, NULL, &numPlatforms );
if( status != CL_SUCCESS )
    fprintf( stderr, "clGetPlatformIDs failed (1)\n" );

fprintf( stderr, "Number of Platforms = %d\n", numPlatforms );

cl_platform_id * platforms = new cl_platform_id[ numPlatforms ];
status = clGetPlatformIDs( numPlatforms, platforms, NULL );
if( status != CL_SUCCESS )
    fprintf( stderr, "clGetPlatformIDs failed (2)\n" );

```



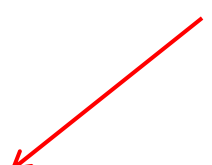
This way of querying information is a recurring OpenCL pattern (get used to it):

	How many to get	Where to put them	How many total there are
status = clGetPlatformIDs(0,	NULL,	&numPlatforms);
status = clGetPlatformIDs(numPlatforms,	platforms,	NULL);

OpenCL Error Codes

14

This one is #define'd as zero.
All the others are negative.



CL_SUCCESS
CL_DEVICE_NOT_FOUND
CL_DEVICE_NOT_AVAILABLE
CL_COMPILER_NOT_AVAILABLE
CL_MEM_OBJECT_ALLOCATION_FAILURE
CL_OUT_OF_RESOURCES
CL_OUT_OF_HOST_MEMORY
CL_PROFILING_INFO_NOT_AVAILABLE
CL_MEM_COPY_OVERLAP
CL_IMAGE_FORMAT_MISMATCH
CL_IMAGE_FORMAT_NOT_SUPPORTED
CL_BUILD_PROGRAM_FAILURE
CL_MAP_FAILURE
CL_INVALID_VALUE
CL_INVALID_DEVICE_TYPE
CL_INVALID_PLATFORM
CL_INVALID_DEVICE
CL_INVALID_CONTEXT

CL_INVALID_QUEUE_PROPERTIES
CL_INVALID_COMMAND_QUEUE
CL_INVALID_HOST_PTR
CL_INVALID_MEM_OBJECT
CL_INVALID_IMAGE_FORMAT_DESCRIPTOR
CL_INVALID_IMAGE_SIZE
CL_INVALID_SAMPLER
CL_INVALID_BINARY
CL_INVALID_BUILD_OPTIONS
CL_INVALID_PROGRAM
CL_INVALID_PROGRAM_EXECUTABLE
CL_INVALID_KERNEL_NAME
CL_INVALID_KERNEL_DEFINITION
CL_INVALID_KERNEL
CL_INVALID_ARG_INDEX
CL_INVALID_ARG_VALUE
CL_INVALID_ARG_SIZE
CL_INVALID_KERNEL_ARGS
CL_INVALID_WORK_DIMENSION



```

struct errorcode
{
    cl_int      statusCode;
    char *      meaning;
}
ErrorCodes[ ] =
{
    { CL_SUCCESS,          " " },
    { CL_DEVICE_NOT_FOUND, "Device Not Found" },
    { CL_DEVICE_NOT_AVAILABLE, "Device Not Available" },
    . . .
    { CL_INVALID_MIP_LEVEL, "Invalid MIP Level" },
    { CL_INVALID_GLOBAL_WORK_SIZE, "Invalid Global Work Size" },
};

void
PrintCLError( cl_int errorCode, char * prefix, FILE *fp )
{
    if( errorCode == CL_SUCCESS )
        return;

    const int numErrorCodes = sizeof( ErrorCodes ) / sizeof( struct errorcode );
    char * meaning = " ";
    for( int i = 0; i < numErrorCodes; i++ )
    {
        if( errorCode == ErrorCodes[i].statusCode )
        {
            meaning = ErrorCodes[i].meaning;
            break;
        }
    }

    fprintf( fp, "%s %s\n", prefix, meaning );
}

```

Querying the Number of Devices on a Platform

16

```
// find out how many devices are attached to each platform and get their ids:  
status = clGetDeviceIDs( platform, CL_DEVICE_TYPE_ALL, 0, NULL, &numDevices );  
devices = new cl_device_id[ numDevices ];  
status = clGetDeviceIDs( platform, CL_DEVICE_TYPE_ALL, numDevices, devices, NULL );
```

Getting Just the GPU Device

```
cl_device_id device;  
status = clGetDeviceIDs( platform, CL_DEVICE_TYPE_GPU, 1, &device, NULL );
```


Querying the Device (this is *really* useful!)

17

```
// find out how many platforms are attached here and get their ids:

cl_uint numPlatforms;
status = clGetPlatformIDs( 0, NULL, &numPlatforms );
if( status != CL_SUCCESS )
    fprintf( stderr, "clGetPlatformIDs failed (1)\n" );

fprintf( OUTPUT, "Number of Platforms = %d\n", numPlatforms );

cl_platform_id *platforms = new cl_platform_id[ numPlatforms ];
status = clGetPlatformIDs( numPlatforms, platforms, NULL );
if( status != CL_SUCCESS )
    fprintf( stderr, "clGetPlatformIDs failed (2)\n" );

cl_uint numDevices;
cl_device_id *devices;

for( int i = 0; i < (int)numPlatforms; i++ )
{
    fprintf( OUTPUT, "Platform #%d:\n", i );
    size_t size;
    char *str;

    clGetPlatformInfo( platforms[i], CL_PLATFORM_NAME, 0, NULL, &size );
    str = new char [ size ];
    clGetPlatformInfo( platforms[i], CL_PLATFORM_NAME, size, str, NULL );
    fprintf( OUTPUT, "\tName   = '%s'\n", str );
    delete[] str;

    clGetPlatformInfo( platforms[i], CL_PLATFORM_VENDOR, 0, NULL, &size );
    str = new char [ size ];
    clGetPlatformInfo( platforms[i], CL_PLATFORM_VENDOR, size, str, NULL );
    fprintf( OUTPUT, "\tVendor = '%s'\n", str );
    delete[] str;
}
```

```

clGetPlatformInfo( platforms[i], CL_PLATFORM_VERSION, 0, NULL, &size );
str = new char [ size ];
clGetPlatformInfo( platforms[i], CL_PLATFORM_VERSION, size, str, NULL );
fprintf( OUTPUT, "\tVersion = '%s'\n", str );
delete[ ] str;

clGetPlatformInfo( platforms[i], CL_PLATFORM_PROFILE, 0, NULL, &size );
str = new char [ size ];
clGetPlatformInfo( platforms[i], CL_PLATFORM_PROFILE, size, str, NULL );
fprintf( OUTPUT, "\tProfile = '%s'\n", str );
delete[ ] str;

// find out how many devices are attached to each platform and get their ids:

status = clGetDeviceIDs( platforms[i], CL_DEVICE_TYPE_ALL, 0, NULL, &numDevices );
if( status != CL_SUCCESS )
    fprintf( stderr, "clGetDeviceIDs failed (2)\n" );

devices = new cl_device_id[ numDevices ];
status = clGetDeviceIDs( platforms[i], CL_DEVICE_TYPE_ALL, numDevices, devices, NULL );
if( status != CL_SUCCESS )
    fprintf( stderr, "clGetDeviceIDs failed (2)\n" );

for( int j = 0; j < (int)numDevices; j++ )
{
    fprintf( OUTPUT, "\tDevice #%d:\n", j );
    size_t size;
    cl_device_type type;
    cl_uint ui;
    size_t sizes[3] = { 0, 0, 0 };

    clGetDeviceInfo( devices[j], CL_DEVICE_TYPE, sizeof(type), &type, NULL );
    fprintf( OUTPUT, "\t\tType = 0x%04x = ", type );

```

```

switch( type )
{
    case CL_DEVICE_TYPE_CPU:
        fprintf( OUTPUT, "CL_DEVICE_TYPE_CPU\n" );
        break;
    case CL_DEVICE_TYPE_GPU:
        fprintf( OUTPUT, "CL_DEVICE_TYPE_GPU\n" );
        break;
    case CL_DEVICE_TYPE_ACCELERATOR:
        fprintf( OUTPUT, "CL_DEVICE_TYPE_ACCELERATOR\n" );
        break;
    default:
        fprintf( OUTPUT, "Other...\n" );
        break;
}

clGetDeviceInfo( devices[j], CL_DEVICE_VENDOR_ID, sizeof(ui), &ui, NULL );
fprintf( OUTPUT, "\t\tDevice Vendor ID = 0x%04x\n", ui );

clGetDeviceInfo( devices[j], CL_DEVICE_MAX_COMPUTE_UNITS, sizeof(ui), &ui, NULL );
fprintf( OUTPUT, "\t\tDevice Maximum Compute Units = %d\n", ui );

clGetDeviceInfo( devices[j], CL_DEVICE_MAX_WORK_ITEM_DIMENSIONS, sizeof(ui), &ui, NULL );
fprintf( OUTPUT, "\t\tDevice Maximum Work Item Dimensions = %d\n", ui );

clGetDeviceInfo( devices[j], CL_DEVICE_MAX_WORK_ITEM_SIZES, sizeof(sizes), sizes, NULL );
fprintf( OUTPUT, "\t\tDevice Maximum Work Item Sizes = %d x %d x %d\n", sizes[0], sizes[1], sizes[2] );

clGetDeviceInfo( devices[j], CL_DEVICE_MAX_WORK_GROUP_SIZE, sizeof(size), &size, NULL );
fprintf( OUTPUT, "\t\tDevice Maximum Work Group Size = %d\n", size );

clGetDeviceInfo( devices[j], CL_DEVICE_MAX_CLOCK_FREQUENCY, sizeof(ui), &ui, NULL );
fprintf( OUTPUT, "\t\tDevice Maximum Clock Frequency = %d MHz\n", ui );

```

Typical Values from Querying the Device

20

Number of Platforms = 1

Platform #0:

 Name = 'NVIDIA CUDA'

 Vendor = 'NVIDIA Corporation'

 Version = 'OpenCL 1.1 CUDA 4.1.1'

 Profile = 'FULL_PROFILE'

 Device #0:

 Type = 0x0004 = CL_DEVICE_TYPE_GPU

 Device Vendor ID = 0x10de

 Device Maximum Compute Units = 15

 Device Maximum Work Item Dimensions = 3

 Device Maximum Work Item Sizes = 1024 x 1024 x 64

 Device Maximum Work Group Size = 1024

 Device Maximum Clock Frequency = 1401 MHz

 Kernel Maximum Work Group Size = 1024

 Kernel Compile Work Group Size = 0 x 0 x 0

 Kernel Local Memory Size = 0



Querying to see what extensions are supported on this device

21

```
size_t extensionSize;

clGetDeviceInfo( device, CL_DEVICE_EXTENSIONS, 0, NULL, &extensionSize );
char *extensions = new char [extensionSize];
clGetDeviceInfo( devices, CL_DEVICE_EXTENSIONS, extensionSize, extensions, NULL );

fprintf( stderr, "\nDevice Extensions:\n" );
for( int i = 0; i < (int)strlen(extensions); i++ )
{
    if( extensions[ i ] == ' ' )
        extensions[ i ] = '\n';
}
fprintf( stderr, "%s\n", extensions );
delete [ ] extensions;
```

Querying to see what extensions are supported on this device

22

This is the big one you are looking for. It shows that this OpenCL system can interoperate with OpenGL.

Device Extensions:

cl_khr_byte_addressable_store
cl_khr_icd
cl_khr_gl_sharing
cl_nv_d3d9_sharing
cl_nv_d3d10_sharing
cl_khr_d3d10_sharing
cl_nv_d3d11_sharing
cl_nv_compiler_options
cl_nv_device_attribute_query
cl_nv_pragma_unroll

This one is handy too. It shows that this OpenCL system can support 64-bit floating point (i.e., double precision).

cl_khr_global_int32_base_atomics
cl_khr_global_int32_extended_atomics
cl_khr_local_int32_base_atomics
cl_khr_local_int32_extended_atomics
cl_khr_fp64

1. Program header
2. Allocate the host memory buffers
3. Create an OpenCL context
4. Create an OpenCL command queue
5. Allocate the device memory buffers
6. Write the data from the host buffers to the device buffers
7. Read the kernel code from a file
8. Compile and link the kernel code
9. Create the kernel object
10. Setup the arguments to the kernel object
11. Enqueue the kernel object for execution
12. Read the results buffer back from the device to the host
13. Clean everything up

1. .cpp Program Header

24

```
#include <stdio.h>
#include <math.h>
#include <string.h>
#include <stdlib.h>
#include <omp.h>    // for timing
#include "cl.h"
```


2. Allocate the Host Memory Buffers

25

```
// allocate the host memory buffers:
```

```
float * hA = new float [ NUM_ELEMENTS ];  
float * hB = new float [ NUM_ELEMENTS ];  
float * hC = new float [ NUM_ELEMENTS ];
```

```
// fill the host memory buffers:
```

```
for( int i = 0; i < NUM_ELEMENTS; i++ )  
{  
    hA[ i ] = hB[ i ] = sqrtf( (float) i );  
}
```

```
// array size in bytes (will need this later):
```

```
size_t dataSize = NUM_ELEMENTS * sizeof( float );
```

```
// opencl function return status:
```

```
cl_int status;                                // test against CL_SUCCESS
```

This could have also been done like this:

```
float hA[ NUM_ELEMENTS ];
```

Global memory and the heap typically have lots more space than the stack. So, you do not want to allocate a large array like this as a local variable.

(Here, it's being done on the heap.)

3. Create an OpenCL Context

26

```
cl_context context = clCreateContext( NULL, 1, &device, NULL, NULL, &status );
```

// create a context:

```
cl_context context = clCreateContext( NULL, 1, &device, NULL, NULL, &status );
```

The diagram illustrates the parameters of the `clCreateContext` function call with arrows pointing to each argument:

- properties**: Points to the first `NULL` argument.
- the device**: Points to the `&device` argument.
- Pass in user data**: Points to the second `NULL` argument.
- one device**: Points to the `1` argument.
- Callback**: Points to the third `NULL` argument.
- returned status**: Points to the `&status` argument.

4. Create an OpenCL Command Queue

27

```
// create a command queue:
```

```
cl_command_queue cmdQueue = clCreateCommandQueue( context, device, 0, &status );
```

the context properties

↓ ↓

```
cl_command_queue cmdQueue = clCreateCommandQueue( context, device, 0, &status );
```

↑ ↑

the device returned status

```
cl_mem dA = clCreateBuffer( context, CL_MEM_READ_ONLY,  dataSize, NULL, &status );
cl_mem dB = clCreateBuffer( context, CL_MEM_READ_ONLY,  dataSize, NULL, &status );
cl_mem dC = clCreateBuffer( context, CL_MEM_WRITE_ONLY, dataSize, NULL, &status );
```

```
cl_mem dA = clCreateBuffer( context, CL_MEM_READ_ONLY,  dataSize, NULL, &status );
```



6. Write the Data from the Host Buffers to the Device Buffers

29

```
// enqueue the 2 commands to write data into the device buffers:
```

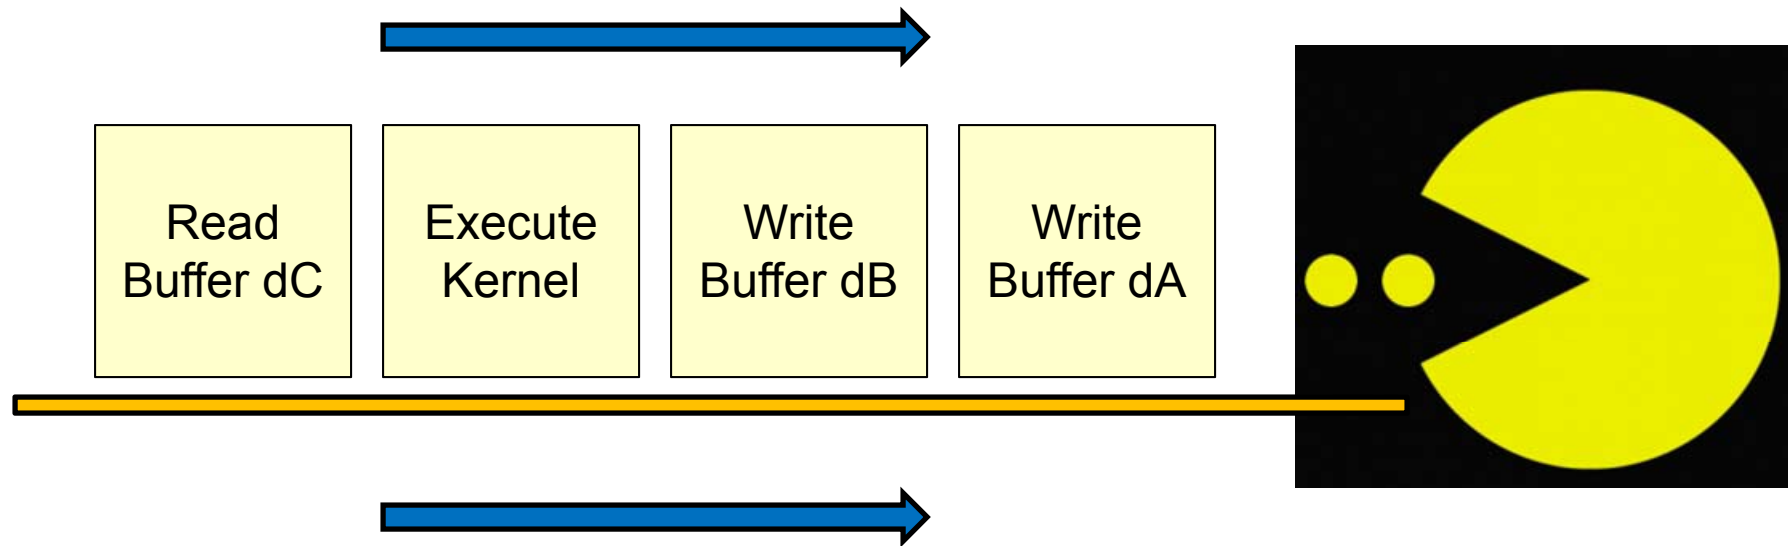
```
status = clEnqueueWriteBuffer( cmdQueue, dA, CL_FALSE, 0, dataSize, hA, 0, NULL, NULL );  
status = clEnqueueWriteBuffer( cmdQueue, dB, CL_FALSE, 0, dataSize, hB, 0, NULL, NULL );
```

The diagram illustrates the parameters of the `clEnqueueWriteBuffer` function call: `status = clEnqueueWriteBuffer(cmdQueue, dA, CL_FALSE, 0, dataSize, hA, 0, NULL, NULL);`. Arrows point from descriptive labels to each parameter:

- command queue** points to `cmdQueue`.
- device buffer** points to `dA`.
- want to block until done?** points to `CL_FALSE`.
- offset** points to `0`.
- # bytes** points to `dataSize`.
- host buffer** points to `hA`.
- # events** points to `0`.
- event object** points to `NULL`.
- event wait list** points to `NULL`.

Enqueueing Works Like a Conveyor Belt

30



```
kernel
void
ArrayMult( global const float *dA, global const float *dB, global float *dC )
{
    int gid = get_global_id( 0 );
    dC[gid] = dA[gid] * dB[gid];
}
```



gid = which element we are
dealing with right now.

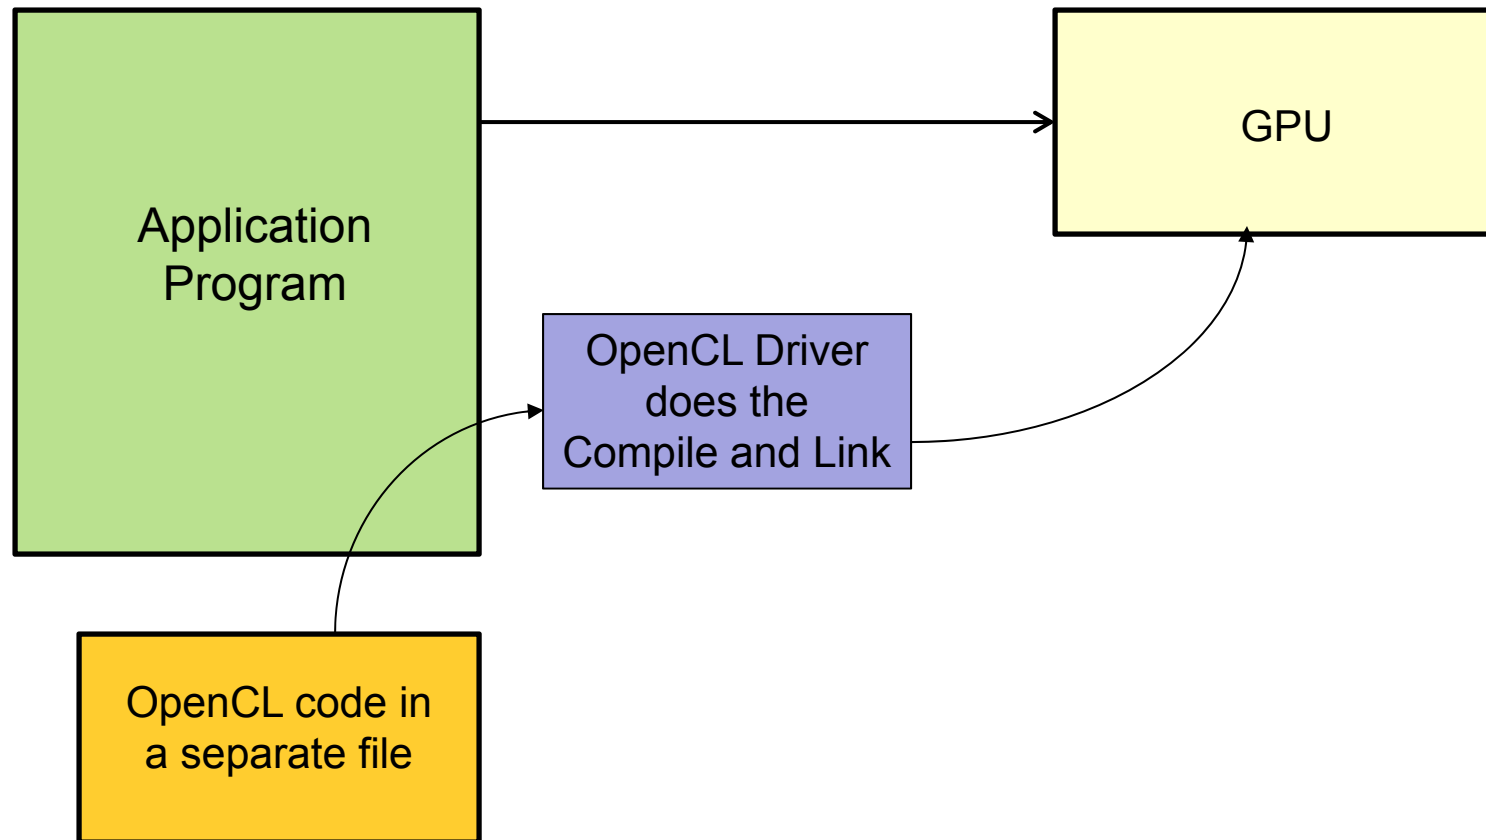
Which dimension's index are we fetching?

0 = X, 1 = Y, 2 = Z

Since this is a 1D problem, X is the only index we need to get.

OpenCL code is compiled in the Driver . . .

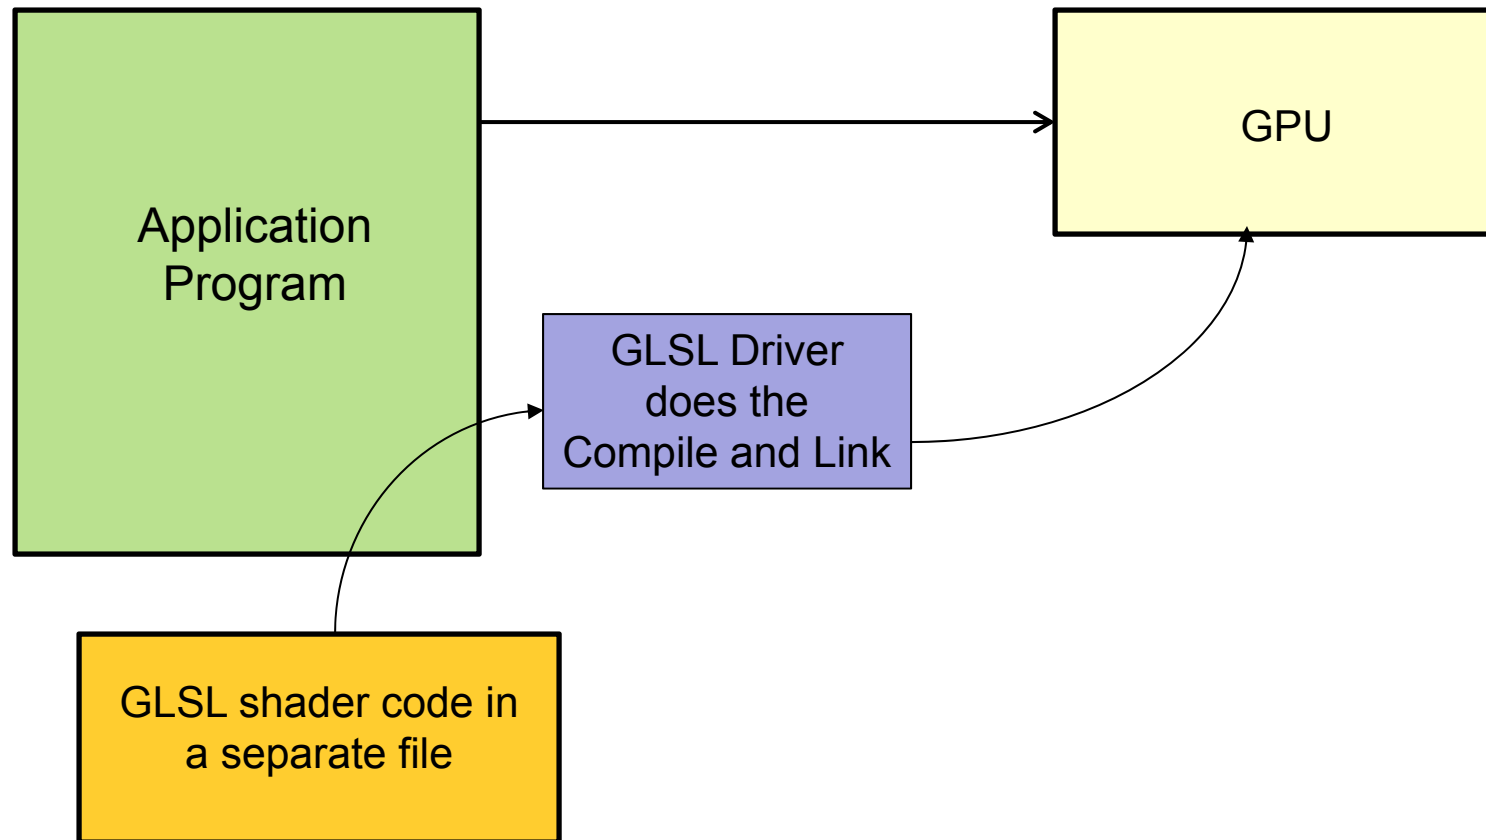
32



```
kernel void
ArrayMult( global float *A, global float *B, global float *C )
{
    int gid = get_global_id ( 0 );

    C[gid] = A[gid] * B[gid];
}
```


(... just like OpenGL's GLSL Shader code is compiled in the driver) 33



```
void main( )  
{  
    vec3 newcolor = texture2D( uTexUnit, vST ).rgb;  
    newcolor = mix( newcolor, vColor.rgb, uBlend );  
    gl_FragColor = vec4(u LightIntensity*newcolor, 1. );  
}
```

7. Read the Kernel Code from a File into a Character Array

34

```
const char *CL_FILE_NAME = { "arraymult.cl" };  
    ...
```

```
FILE *fp = fopen( CL_FILE_NAME, "r" );  
if( fp == NULL )  
{  
    fprintf( stderr, "Cannot open OpenCL source file '%s'\n", CL_FILE_NAME );  
    return 1;  
}
```

```
// read the characters from the opencl kernel program:
```

```
fseek( fp, 0, SEEK_END );  
size_t fileSize = ftell( fp );  
fseek( fp, 0, SEEK_SET );  
char *clProgramText = new char[ fileSize+1 ];  
size_t n = fread( clProgramText, 1, fileSize, fp );  
clProgramText[fileSize] = '\0';  
fclose( fp );
```

"r" should work, since the .cl file is pure ASCII text, but some people report that it doesn't work unless you use "rb"

Watch out for the '\r' + '\n' problem! (See the next slide.)

Some of you will end up having strange, unexplainable problems with your `cs`h scripts, `.cpp` programs, or `.cl` programs. This could be because you are typing your code in on Windows (using Notepad or Wordpad or Word) and then running it on Linux. Windows likes to insert an extra carriage return (`\r`) at the end of each line, which Linux interprets as a garbage character.

You can test this by typing the Linux command:

`od -c loop.csh`

which will show you all the characters, even the `\r` (which you don't want) and the `\n` (newlines, which you do want).

To get rid of the carriage returns, enter the Linux command:

`tr -d '\r' < loop.csh > loop1.csh`

Then run `loop1.csh`

Or, on some systems, there is a utility called *dos2unix* which does this for you:

`dos2unix < loop.csh > loop1.csh`

Sorry about this. Unfortunately, this is a fact of life when you mix Windows and Linux.

8. Compile and Link the Kernel Code

36

```
// create the kernel program on the device:

char * strings [ 1 ];           // an array of strings
strings[0] = clProgramText;
cl_program program = clCreateProgramWithSource( context, 1, (const char **)strings, NULL, &status );
delete [ ] clProgramText;

// build the kernel program on the device:

char *options = { "" };
status = clBuildProgram( program, 1, &device, options, NULL, NULL );
if( status != CL_SUCCESS )
{
    // retrieve and print the error messages:
    size_t size;
    clGetProgramBuildInfo( program, devices[0], CL_PROGRAM_BUILD_LOG, 0, NULL, &size );
    cl_char *log = new cl_char[ size ];
    clGetProgramBuildInfo( program, devices[0], CL_PROGRAM_BUILD_LOG, size, log, NULL );
    fprintf( stderr, "clBuildProgram failed:\n%s\n", log );
    delete [ ] log;
}
```

How does that array-of-strings thing actually work?

37

```
char *ArrayOfStrings[3];
ArrayOfStrings[0] = "...one commonly-used function...";
ArrayOfStrings[1] = "... another commonly-used function. . . ";
ArrayOfStrings[2] = "... the real OpenCL code . . . ";
cl_program program = clCreateProgramWithSource( context, 1, (const char **) ArrayOfStrings, NULL, &status );
```

These are two ways to provide a *single* character buffer:

```
char *buffer[1];
buffer[0] = "... the entire OpenCL code . . . ";
cl_program program = clCreateProgramWithSource( context, 1, (const char **) buffer, NULL, &status );
```

```
char *buffer = "... the entire OpenCL code . . . ";
cl_program program = clCreateProgramWithSource( context, 1, (const char **) &buffer, NULL, &status );
```

Why use an array of strings to hold the OpenCL program, instead of just a single string?

1. You can use the same OpenCL source and insert the appropriate “#defines” at the beginning
2. You can insert a common header file (\approx a .h file)
3. You can simulate a “#include” to re-use common pieces of code

9. Create the Kernel Object

39

```
cl_kernel kernel = clCreateKernel( program, "ArrayMult", &status );
```



10. Setup the Arguments to the Kernel Object

40

```
status = clSetKernelArg( kernel, 0, sizeof(cl_mem), &dA );  
status = clSetKernelArg( kernel, 1, sizeof(cl_mem), &dB );  
status = clSetKernelArg( kernel, 2, sizeof(cl_mem), &dC );
```

kernel
void

ArrayMult(global const float *dA, global const float *dB, global float *dC)

11. Enqueue the Kernel Object for Execution

41

```
size_t globalWorkSize[ 3 ] = { NUM_ELEMENT, 1, 1 };
size_t localWorkSize[ 3 ] = { LOCAL_SIZE, 1, 1 };

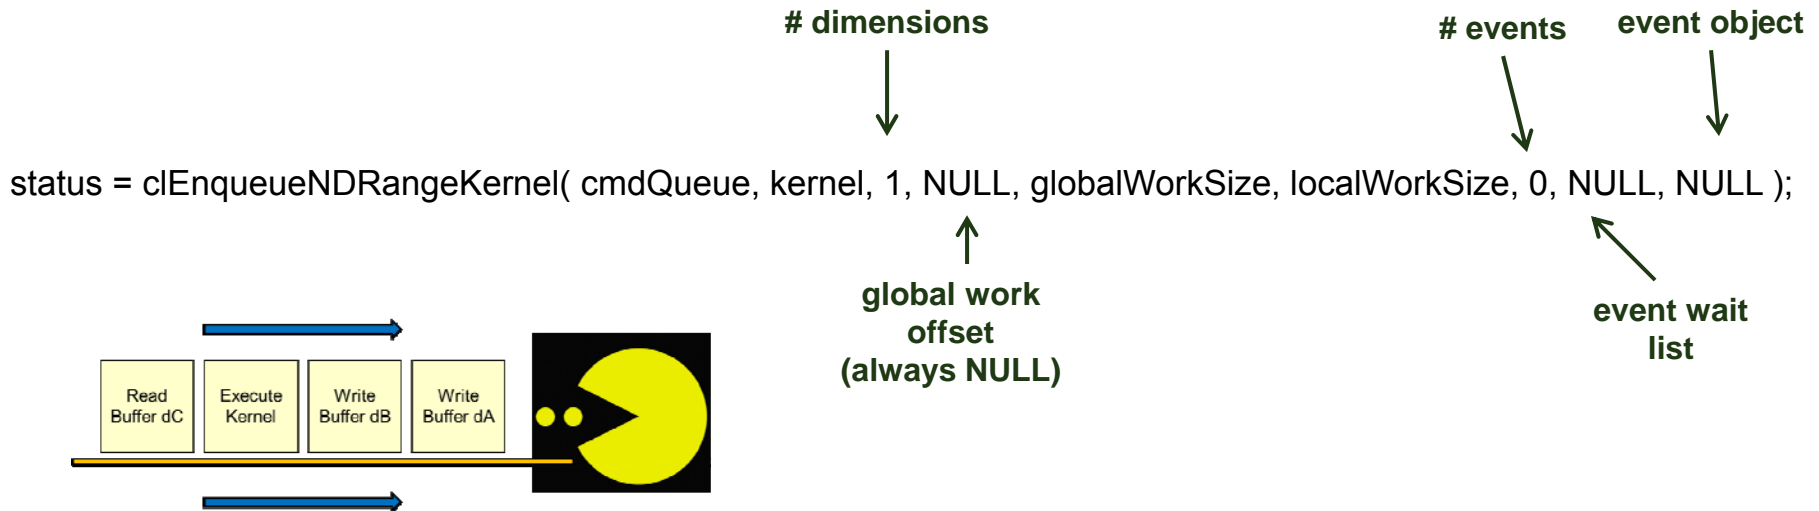
status = clEnqueueBarrier( cmdQueue );

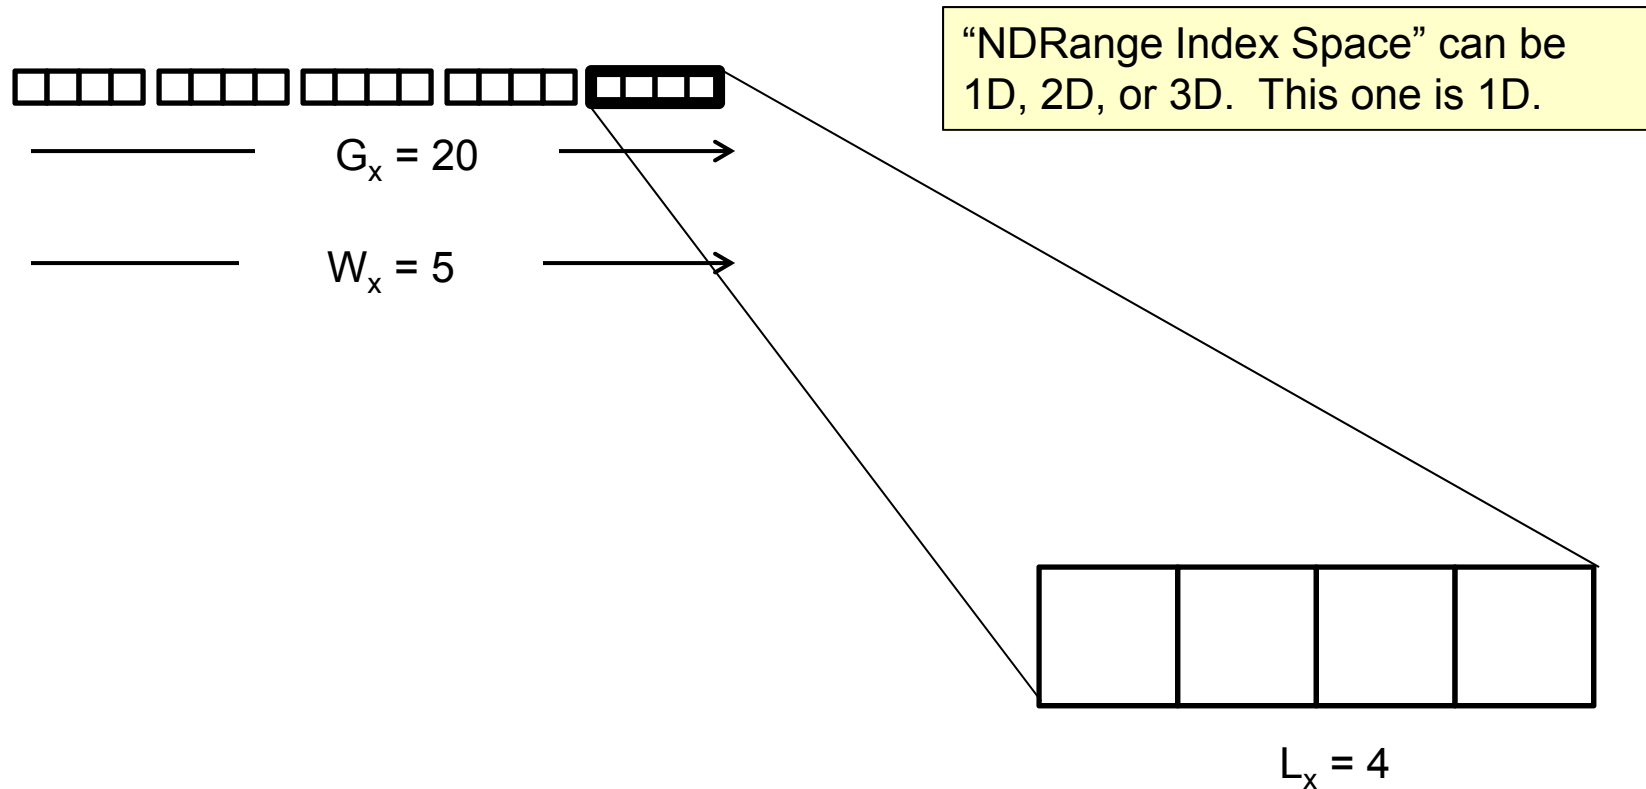
double time0 = omp_get_wtime( );

status = clEnqueueNDRangeKernel( cmdQueue, kernel, 1, NULL, globalWorkSize, localWorkSize, 0, NULL, NULL );

status = clEnqueueBarrier( cmdQueue );

double time1 = omp_get_wtime( );
```






$$\#WorkGroups = \frac{GlobalIndexSpaceSize}{WorkGroupSize}$$


$$5 \times 4 = \frac{20}{4}$$

Work-Groups, Local IDs, and Global IDs

43



“NDRange Index Space” can be 1D, 2D, or 3D. This one is 2D.




$$\#WorkGroups = \frac{GlobalIndexSpaceSize}{WorkGroupSize}$$

$$5 \times 4 = \frac{20 \times 12}{4 \times 3}$$

Work-Groups, Local IDs, and Global IDs

44

“NDRange Index Space” can be 1D, 2D, or 3D. This one is 3D.



Figuring Out What Thread You Are and What Your Thread Environment is Like

```
uint    get_work_dim( ) ;
```

```
size_t  get_global_size( uint dimindx ) ;
```

```
size_t  get_global_id( uint dimindx ) ;
```

```
size_t  get_local_size( uint dimindx ) ;
```

```
size_t  get_local_id( uint dimindx ) ;
```

```
size_t  get_num_groups( uint dimindx ) ;
```

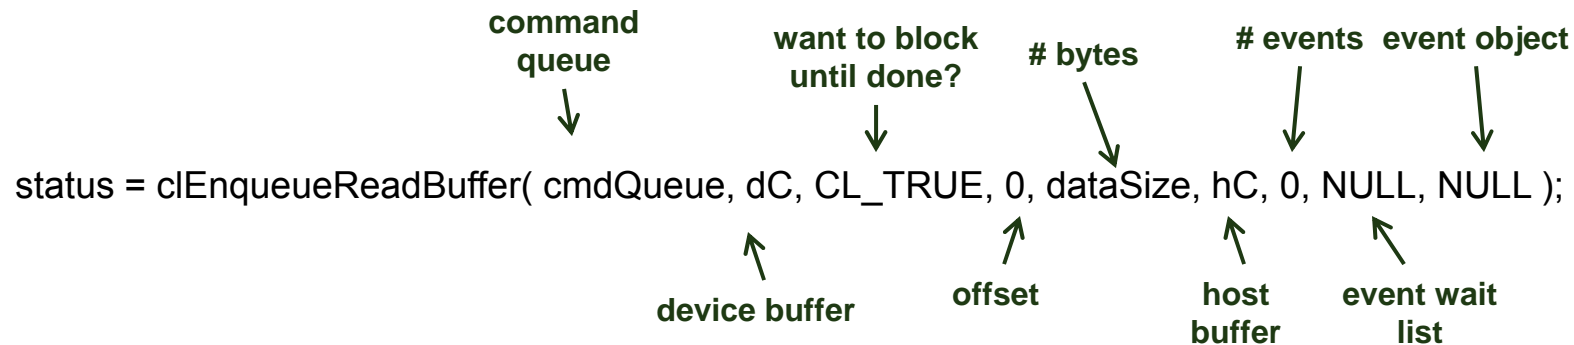
```
size_t  get_group_id( uint dimindx ) ;
```

```
size_t  get_global_offset( uint dimindx ) ;
```

$$0 \leq \textit{dimindx} \leq 2$$

12. Read the Results Buffer Back from the Device to the Host

```
status = clEnqueueReadBuffer( cmdQueue, dC, CL_TRUE, 0, dataSize, hC, 0, NULL, NULL );
```




13. Clean Everything Up


47

```
// clean everything up:  
  
clReleaseKernel(      kernel );  
clReleaseProgram(     program );  
clReleaseCommandQueue( cmdQueue );  
clReleaseMemObject(   dA );  
clReleaseMemObject(   dB );  
clReleaseMemObject(   dC );  
  
delete [ ] hA;  
delete [ ] hB;  
delete [ ] hC;
```

Array Multiplication Performance: What is a Good Work-Group Size?



Array Multiplication Performance: What is a Good Work-Group Size?



Writing the .cl Program's Binary Code

```
size_t binary_sizes;  
status = clGetProgramInfo( Program, CL_PROGRAM_BINARY_SIZES, 0, NULL, &binary_sizes );  
  
size_t size;  
status = clGetProgramInfo( Program, CL_PROGRAM_BINARY_SIZES, sizeof(size_t), &size, NULL );  
  
unsigned char *binary = new unsigned char [ size ];  
status = clGetProgramInfo( Program, CL_PROGRAM_BINARIES, size, &binary, NULL );  
  
FILE *fpbin = fopen( "particles.nv", "wb" );  
if( fpbin == NULL )  
{  
    fprintf( stderr, "Cannot create 'particles.bin'\n" );  
}  
else  
{  
    fwrite( binary, 1, size, fpbin );  
    fclose( fpbin );  
}  
delete [ ] binary;
```

Importing that Binary Code back In:

8. Compile and Link the Kernel Code

Instead of doing this:

```
char * strings [ 1 ];
strings[0] = clProgramText;
cl_program program = clCreateProgramWithSource( context, 1, (const char **)strings, NULL, &status );
delete [ ] clProgramText;
```

You would do this:


```
unsigned char byteArray[ numBytes ];
cl_program program = clCreateProgramWithBinary( context, 1, &device, &numBytes, &byteArray, &binaryStatus, &status );
delete [ ] byteArray;
```

And you still have to do this:

```
char *options = { "" };
status = clBuildProgram( program, 1, &device, options, NULL, NULL );
if( status != CL_SUCCESS )
{
    size_t size;
    clGetProgramBuildInfo( program, device, CL_PROGRAM_BUILD_LOG, 0, NULL, &size );
    cl_char *log = new cl_char[ size ];
    clGetProgramBuildInfo( program, device, CL_PROGRAM_BUILD_LOG, size, log, NULL );
    fprintf( stderr, "clBuildProgram failed:\n%s\n", log );
    delete [ ] log;
}
```

Billion OpenCL Multiplies per Second on *rabbit's* NVIDIA Titan Black

(Array size = 64M)



Billion OpenCL Multiplies per Second on *rabbit's* NVIDIA Titan Black

(Local Size = 64)

